

# Accelerating Genome Sequence Analysis via Efficient Hardware/Algorithm Co-Design

**Damlasenol Cali, Ph.D.**

<https://damlasenolcali.github.io/>

[damlasenolcali@gmail.com](mailto:damlasenolcali@gmail.com)

Staff Software Engineer, Hardware Acceleration  
Bionano Genomics

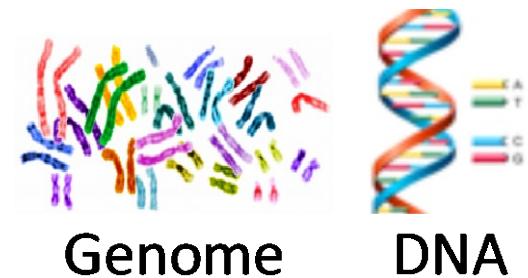
**AACBB Workshop @ ISCA 2022**  
June 18, 2022

# Genome Sequencing

---

- **Genome sequencing:** Enables us to determine the order of the DNA sequence in an organism's genome

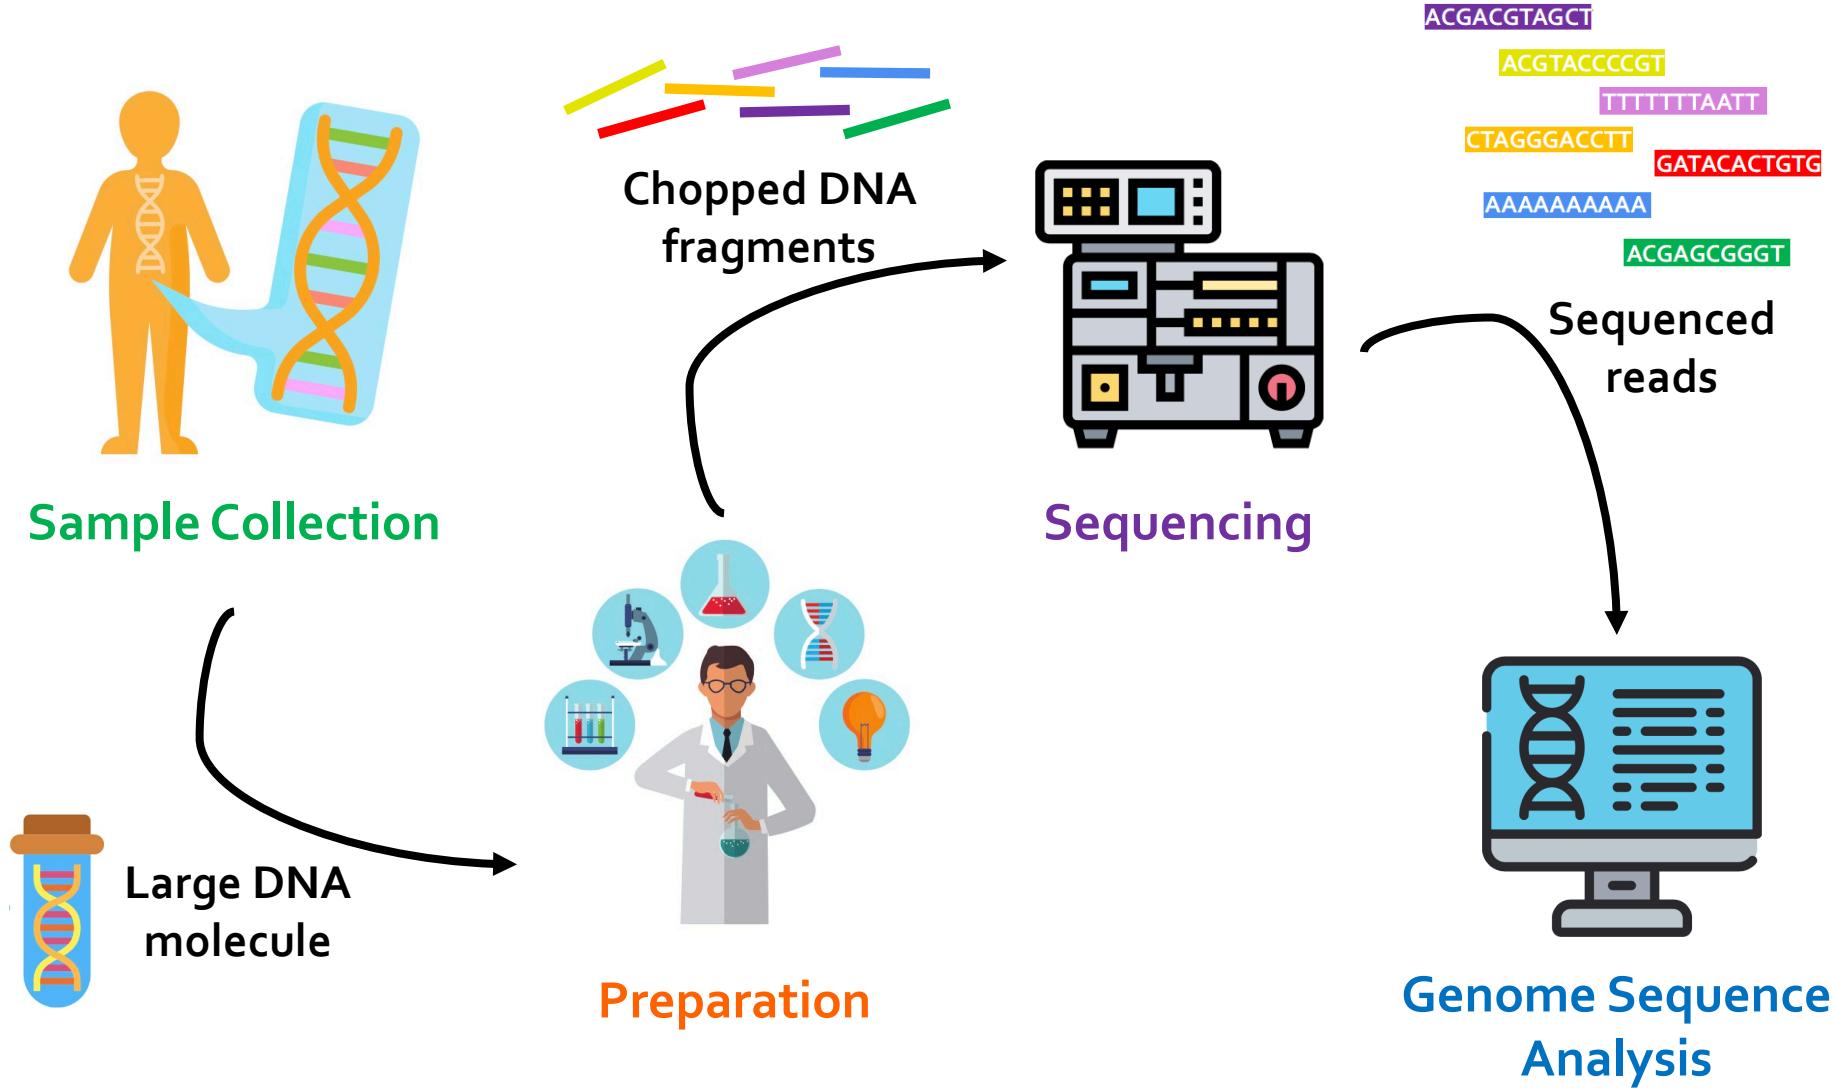
- Plays a **pivotal role** in:
  - Personalized medicine
  - Outbreak tracing
  - Understanding of evolution



- **Challenges:**

- There is no sequencing machine that takes long DNA as an input, and gives the complete sequence as output
- Sequencing machines extract **small randomized fragments** of the original DNA sequence

# Genome Sequencing (cont'd.)



# Sequencing Technologies



Oxford Nanopore  
(ONT)



PacBio

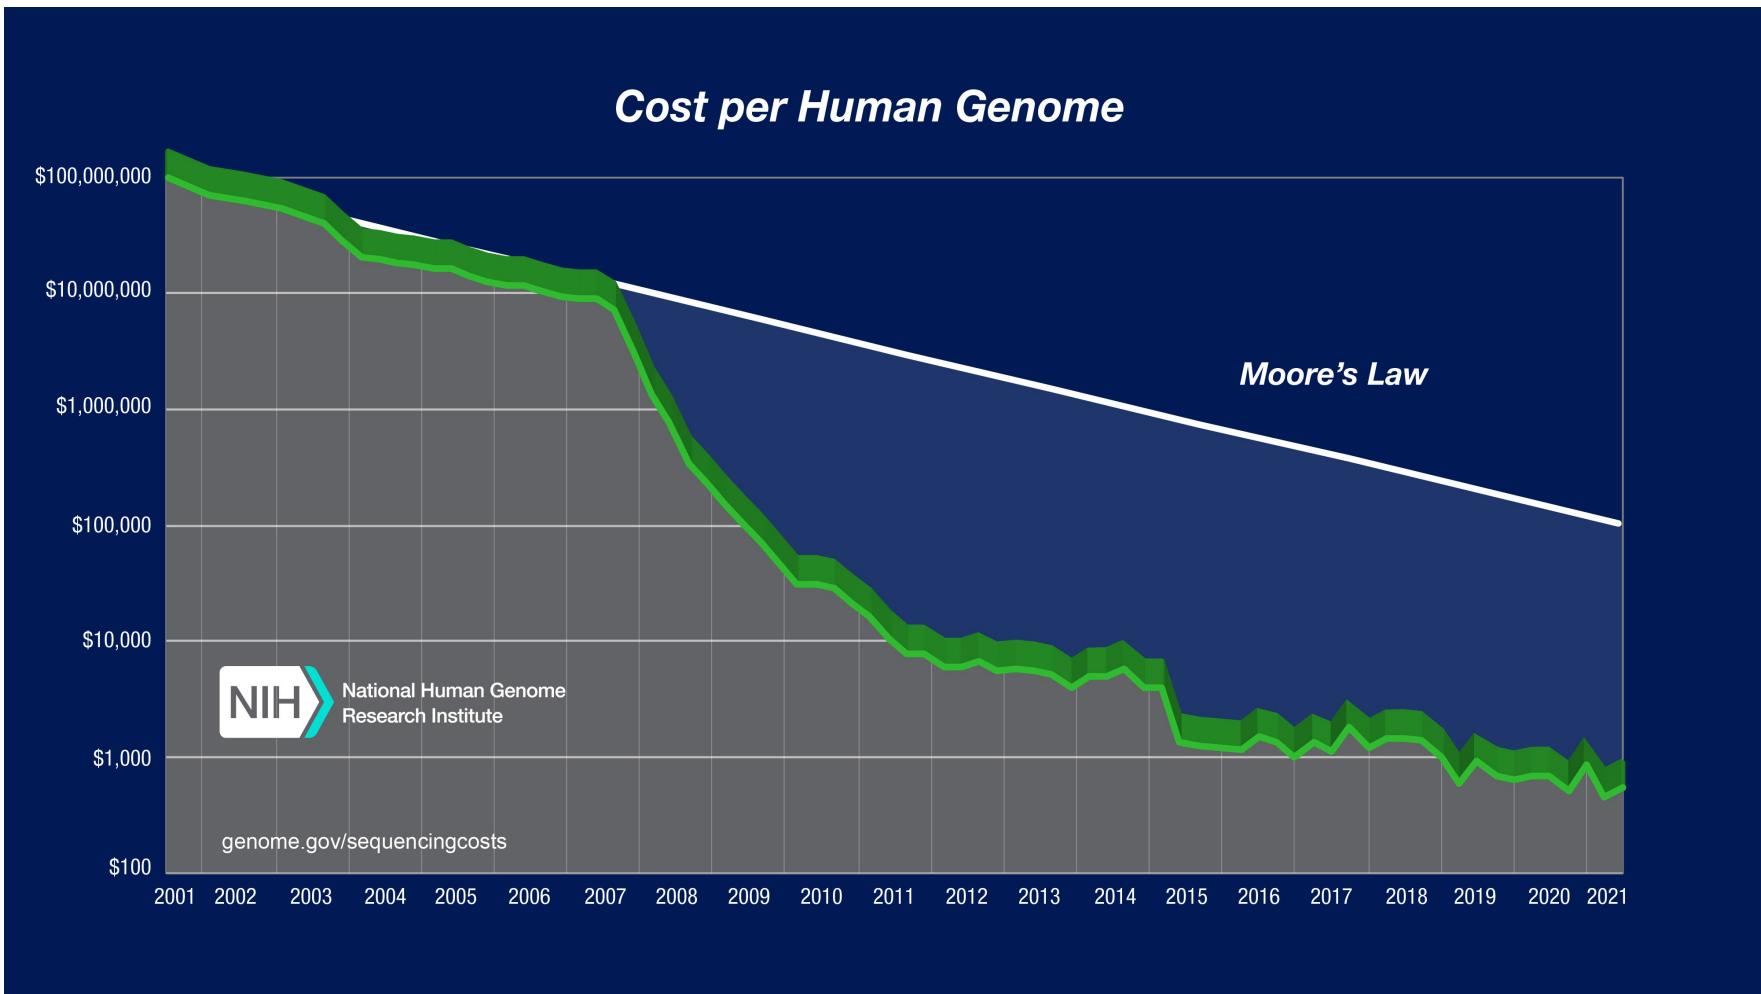
Illumina



***Short reads:* a few hundred base pairs and error rate of ~0.1%**

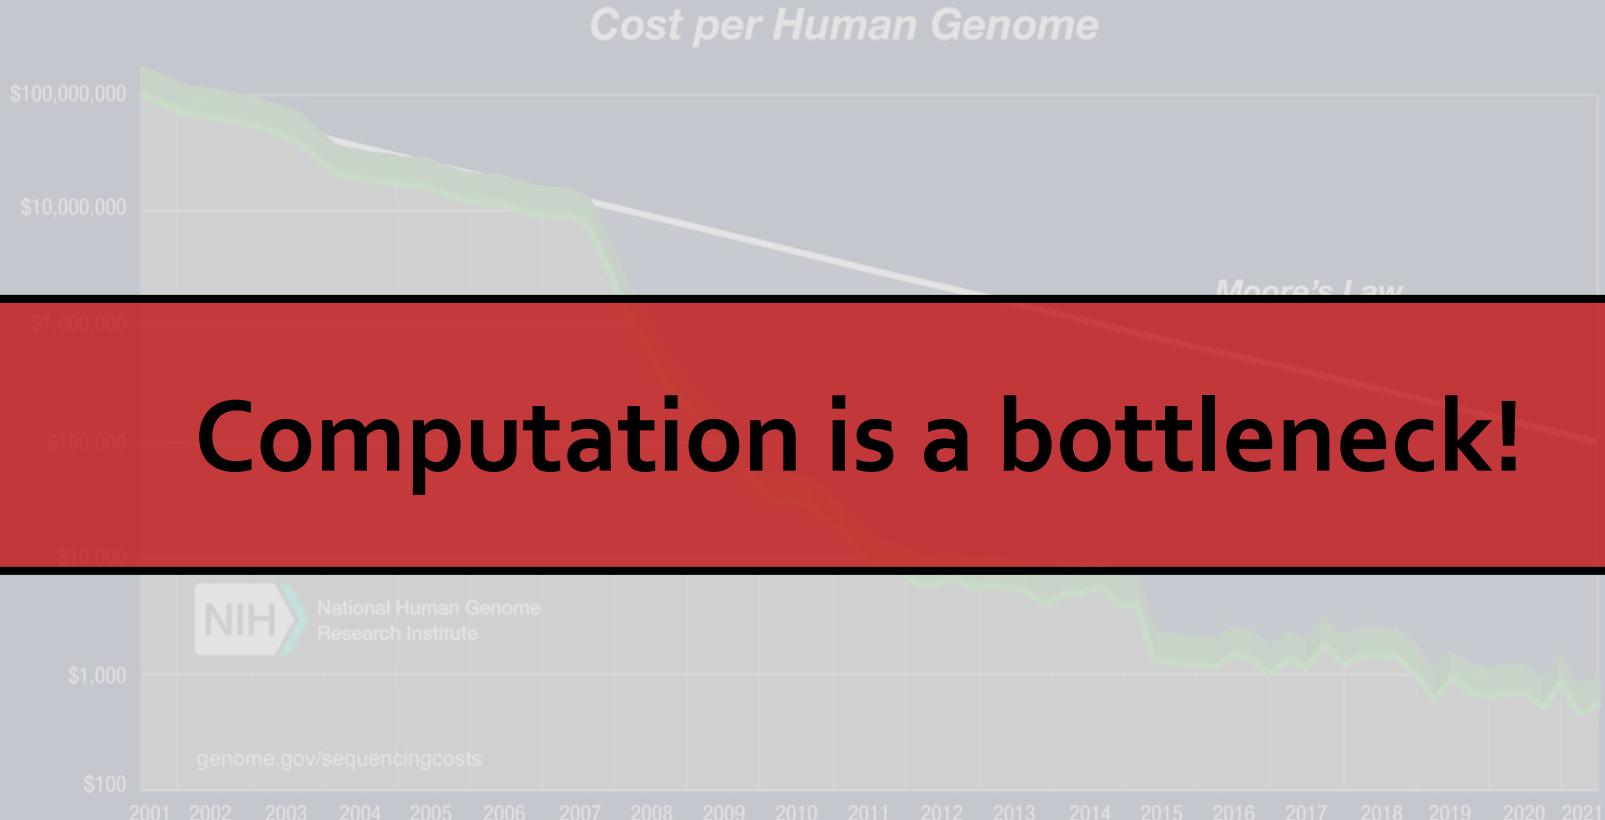
***Long reads:* thousands to millions of base pairs and error rate of 5–10%**

# Current State of Sequencing



\*From NIH (<https://www.genome.gov/about-genomics/fact-sheets/DNA-Sequencing-Costs-Data>)

# Current State of Sequencing (cont'd.)



\*From NIH (<https://www.genome.gov/about-genomics/fact-sheets/DNA-Sequencing-Costs-Data>)

# Our Goal & Approach

---

- Our Goal:  
Accelerating genome sequence analysis by **efficient hardware/algorithm co-design**
  
- Our Approach:
  - (1) Analyze the **multiple steps** and the **associated tools** in the genome sequence analysis pipeline,
  - (2) Expose the **tradeoffs** between accuracy, performance, memory usage and scalability, and
  - (3) Co-design **fast and efficient algorithms** along with **scalable and energy-efficient customized hardware accelerators** for the key bottleneck steps of the pipeline

# Research Contributions

Bottleneck analysis of genome assembly pipeline for long reads

*[Briefings in Bioinformatics, 2018]*

**GenASM:** Approximate string matching framework for genome sequence analysis

*[MICRO 2020]*

BitMAC: FPGA-based near-memory acceleration of bitvector-based sequence alignment

*[Ongoing]*

**SeGram:** Universal genomic mapping accelerator for both sequence-to-graph and sequence-to-sequence mapping

*[ISCA 2022]*

# Nanopore Sequencing & Tools [BiB 2018]

---

Damla Senol Cali, Jeremie S. Kim, Saugata Ghose, Can Alkan, and Onur Mutlu,  
"Nanopore Sequencing Technology and Tools for Genome Assembly:  
Computational Analysis of the Current State, Bottlenecks and Future  
Directions"

*Briefings in Bioinformatics*, April 2018.

## **Nanopore Sequencing Technology and Tools for Genome Assembly: Computational Analysis of the Current State, Bottlenecks and Future Directions**

**Damla Senol Cali<sup>1,\*</sup>, Jeremie S. Kim<sup>1,3</sup>, Saugata Ghose<sup>1</sup>, Can Alkan<sup>2\*</sup>  
and Onur Mutlu<sup>3,1\*</sup>**

<sup>1</sup>Department of Electrical and Computer Engineering, Carnegie Mellon University, Pittsburgh, PA, USA

<sup>2</sup>Department of Computer Engineering, Bilkent University, Bilkent, Ankara, Turkey

<sup>3</sup>Department of Computer Science, Systems Group, ETH Zürich, Zürich, Switzerland

# Key Findings

**Goal 1:**

**High-performance and low-power**

**Goal 2:**

**Memory-efficient**

**Goal 3:**

**Scalable/highly-parallel**

# Research Contributions

Bottleneck analysis of genome assembly pipeline for long reads

*[Briefings in Bioinformatics, 2018]*

**GenASM:** Approximate string matching framework for genome sequence analysis

*[MICRO 2020]*

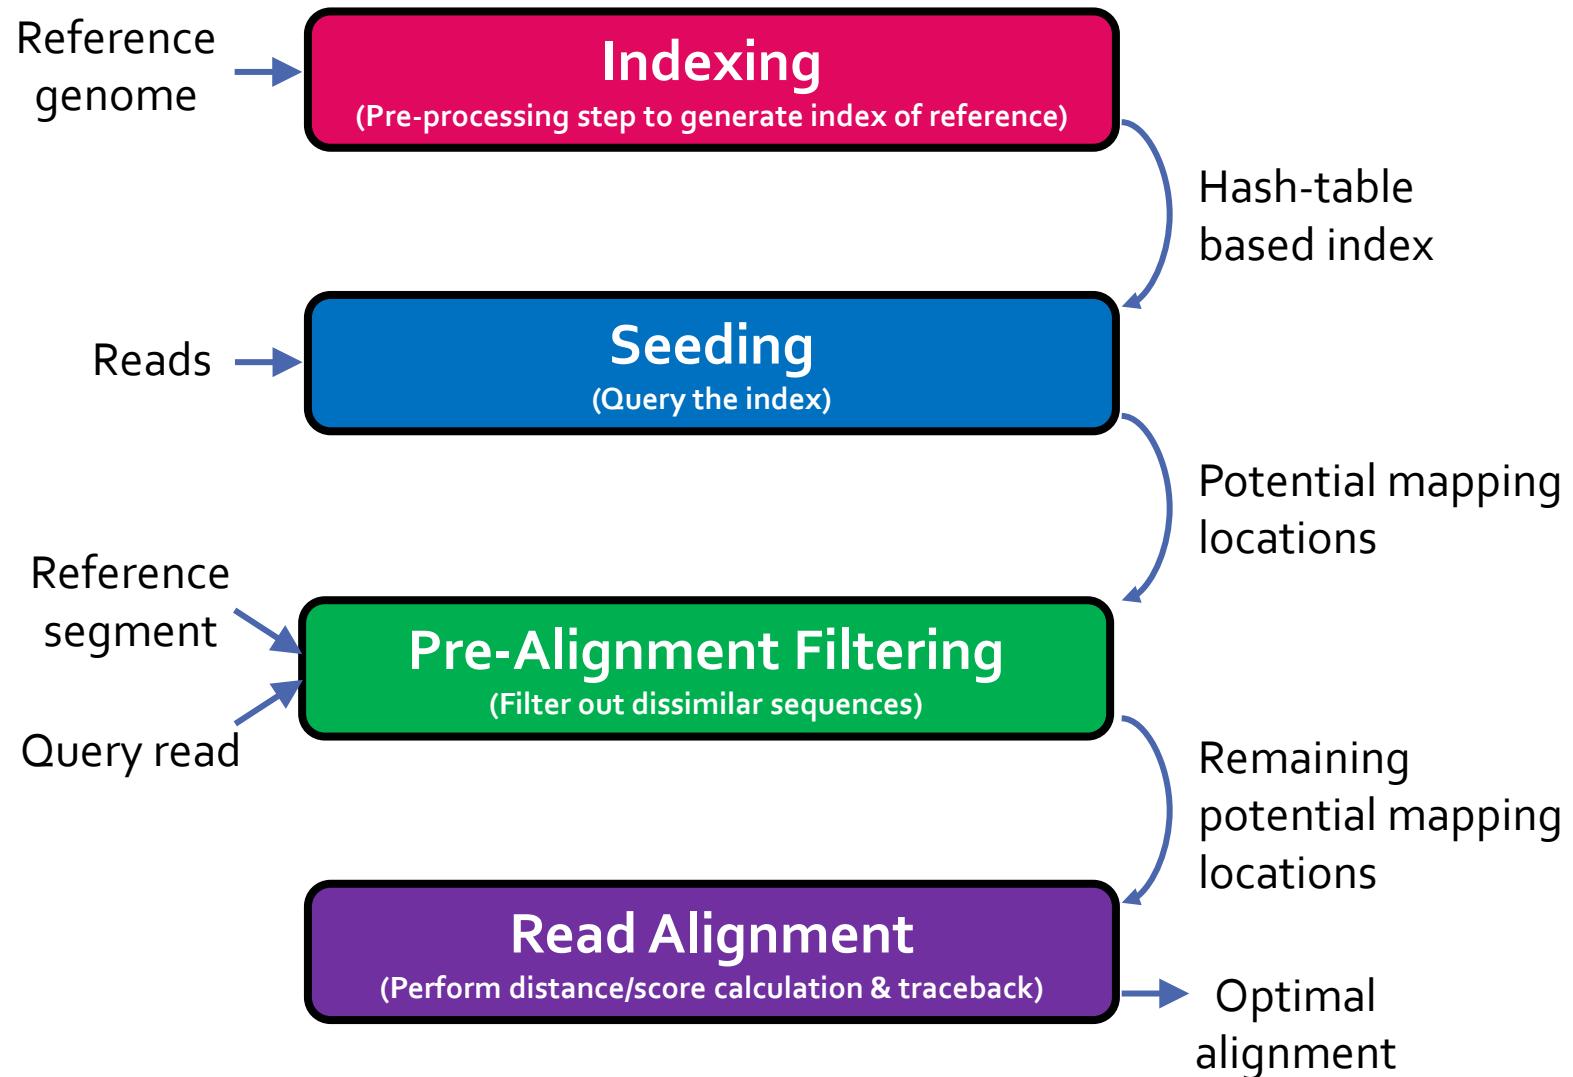
BitMAC: FPGA-based near-memory acceleration of bitvector-based sequence alignment

*[Ongoing]*

SeGram: Universal genomic mapping accelerator for both sequence-to-graph and sequence-to-sequence mapping

*[ISCA 2022]*

# Read Mapping Pipeline



# GSA with Read Mapping

---

- ❑ **Read mapping:** *First key step* in genome sequence analysis (GSA)
  - Aligns **reads** to one or more possible locations within the **reference genome**, and
  - Finds the **matches** and **differences** between the read and the reference genome segment at that location
- ❑ Multiple steps of read mapping require **approximate string matching**
  - Approximate string matching (ASM) enables read mapping to account for **sequencing errors** and **genetic variations** in the reads
- ❑ Bottlenecked by the **computational power** and **memory bandwidth** limitations of existing systems

# Approximate String Matching

---

- Sequenced genome **may not exactly map** to the reference genome due to **genetic variations** and **sequencing errors**

*Reference:* AAA T G T T T A T G C T A C T G  
*Read:* AAA T G T T T A T G C T A C T G

*deletion*      *substitution*      *insertion*

- **Approximate string matching (ASM):**
  - Detect the **differences** and **similarities** between two sequences
  - In genomics, ASM is required to:
    - Find the **minimum edit distance** (i.e., total number of differences)
    - Find the **optimal alignment** with a **traceback** step
      - Sequence of matches, substitutions, insertions and deletions, along with their positions
  - Usually implemented as a **dynamic programming (DP) based algorithm**

# Bitap Algorithm

---

- Bitap<sup>1,2</sup> performs ASM with fast and simple bitwise operations
  - Amenable to efficient hardware acceleration
  - Computes the **minimum edit distance** between a **text** (e.g., reference genome) and a **pattern** (e.g., read) with a maximum of  $k$  errors
- **Step 1: Pre-processing (per pattern)**
  - Generate a **pattern bitmask (PM)** for each character in the alphabet (A, C, G, T)
  - Each PM indicates if character exists at each position of the pattern
- **Step 2: Searching (Edit Distance Calculation)**
  - Compare all characters of the **text** with the **pattern** by using:
    - Pattern bitmasks
    - Status bitvectors that hold the partial matches
    - Bitwise operations

[1] R. A. Baeza-Yates and G. H. Gonnet. "A New Approach to Text Searching." *CACM*, 1992.  
[2] S. Wu and U. Manber. "Fast Text Searching: Allowing Errors." *CACM*, 1992.

# Limitations of Bitap

---

## 1) Data Dependency Between Iterations:

- Two-level data dependency forces the consecutive iterations to take place sequentially

# Bitap Algorithm (cont'd.)

## □ Step 2: Edit Distance Calculation

For each character of the text (char):

Copy previous R bitvectors as oldR

$R[0] = (\text{oldR}[0] \ll 1) | \text{PM}[\text{char}]$

For  $d = 1 \dots k$ :

deletion =  $\text{oldR}[d-1]$

substitution =  $\text{oldR}[d-1] \ll 1$

insertion =  $R[d-1] \ll 1$

match =  $(\text{oldR}[d] \ll 1) | \text{PM}[\text{char}]$

$R[d] = \text{deletion} \& \text{mismatch} \& \text{insertion} \& \text{match}$

Check MSB of  $R[d]$ :

If 1, no match.

If 0, match with  $d$  many errors.

Large number of iterations

# Bitap Algorithm (cont'd.)

## □ Step 2: Edit Distance Calculation

For each character of the text (char):

Copy previous R bitvectors as oldR

$R[0] = (\text{oldR}[0] \ll 1) | \text{PM}[\text{char}]$

For  $d = 1 \dots k$ :

deletion       $= \text{oldR}[d-1]$

substitution     $= \text{oldR}[d-1] \ll 1$

insertion       $= R[d-1] \ll 1$

match             $= (\text{oldR}[d] \ll 1) | \text{PM}[\text{char}]$

$R[d] = \text{deletion} \& \text{mismatch} \& \text{insertion} \& \text{match}$

Check MSB of  $R[d]$ :

If 1, no match.

If 0, match with  $d$  many errors.

Data dependency  
between iterations  
(i.e., no  
parallelization)

# Limitations of Bitap

---

## 1) Data Dependency Between Iterations:

- Two-level data dependency forces the consecutive iterations to take place sequentially

## 2) No Support for Traceback:

- Bitap does not include any support for optimal alignment identification

# Bitap Algorithm (cont'd.)

## □ Step 2: Edit Distance Calculation

For each character of the text (char):

Copy previous R bitvectors as oldR

$R[0] = (\text{oldR}[0] \ll 1) | \text{PM}[\text{char}]$

For  $d = 1 \dots k$ :

deletion	$= \text{oldR}[d-1]$
substitution	$= \text{oldR}[d-1] \ll 1$
insertion	$= R[d-1] \ll 1$
match	$= (\text{oldR}[d] \ll 1)   \text{PM}[\text{char}]$

Does *not* store and process these intermediate bitvectors to find the optimal alignment (i.e., no traceback)

$R[d] = \text{deletion} \& \text{mismatch} \& \text{insertion} \& \text{match}$

Check MSB of  $R[d]$ :

If 1, no match.

If 0, match with  $d$  many errors.

# Limitations of Bitap

## 1) Data Dependency Between Iterations:

- Two-level data dependency forces the consecutive iterations to take place sequentially

## 2) No Support for Traceback:

- Bitap does not include any support for optimal alignment identification

## 3) No Support for Long Reads:

- Each bitvector has a length equal to the length of the pattern
- Bitwise operations are performed on these bitvectors

## 4) Limited Compute Parallelism:

- Text-level parallelism
- Limited by the number of compute units in existing systems

## 5) Limited Memory Bandwidth:

- High memory bandwidth required to read and write the computed bitvectors to memory

Algorithm

Hardware

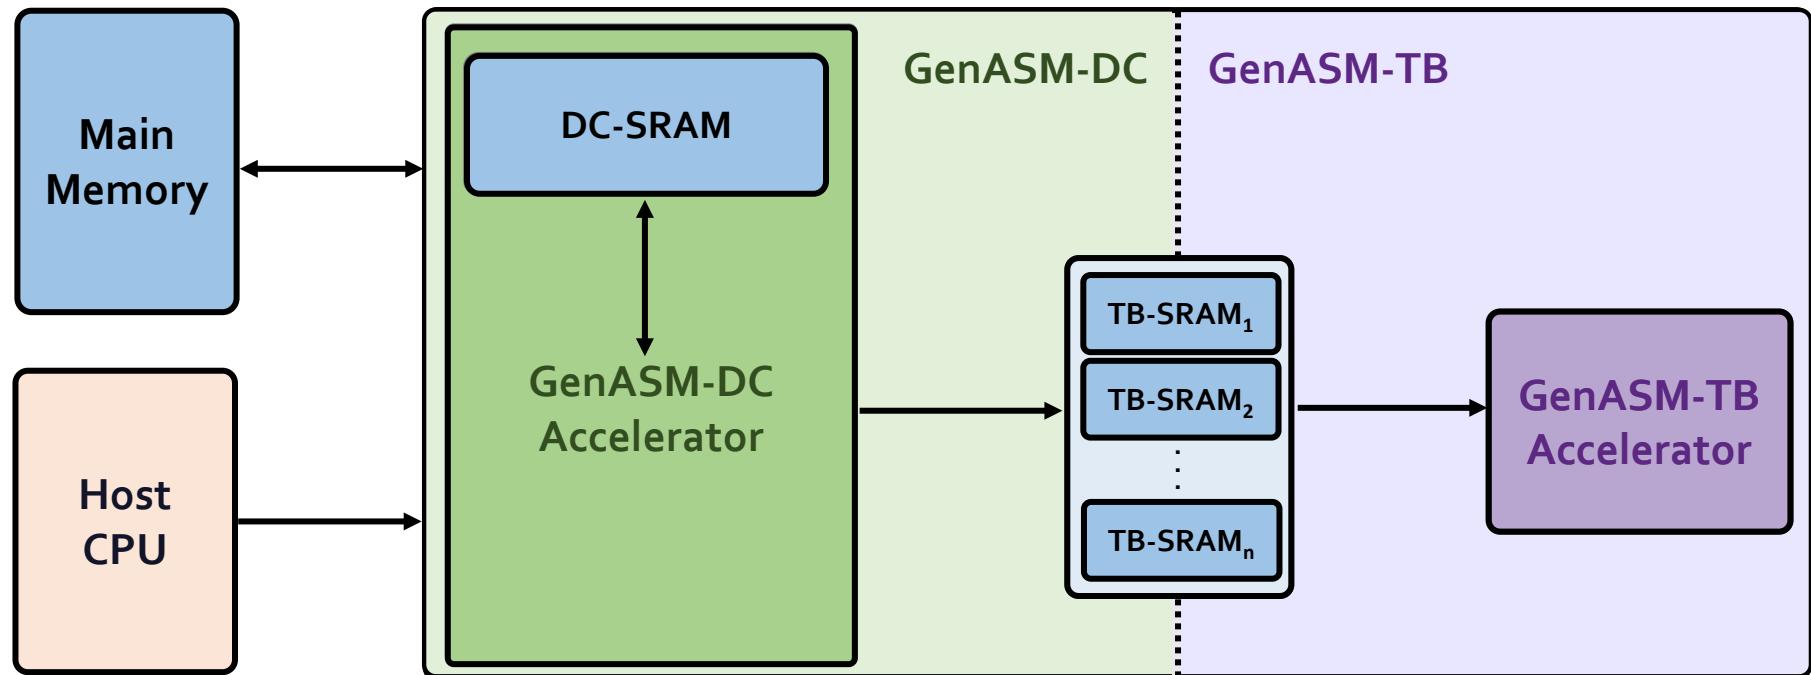
# GenASM: ASM Framework for GSA

## Our Goal:

Accelerate approximate string matching  
by designing a fast and flexible framework,  
which can accelerate *multiple steps* of genome sequence analysis

- ❑ GenASM: First ASM acceleration framework for GSA
  - Approximate string matching (ASM) acceleration framework based on the Bitap algorithm
- ❑ We overcome the five limitations that hinder Bitap's use in GSA:
  - Modified and extended ASM algorithm
    - Highly-parallel Bitap with long read support
    - Novel bitvector-based algorithm to perform *traceback*
  - Specialized, low-power and area-efficient hardware for both modified Bitap and novel traceback algorithms

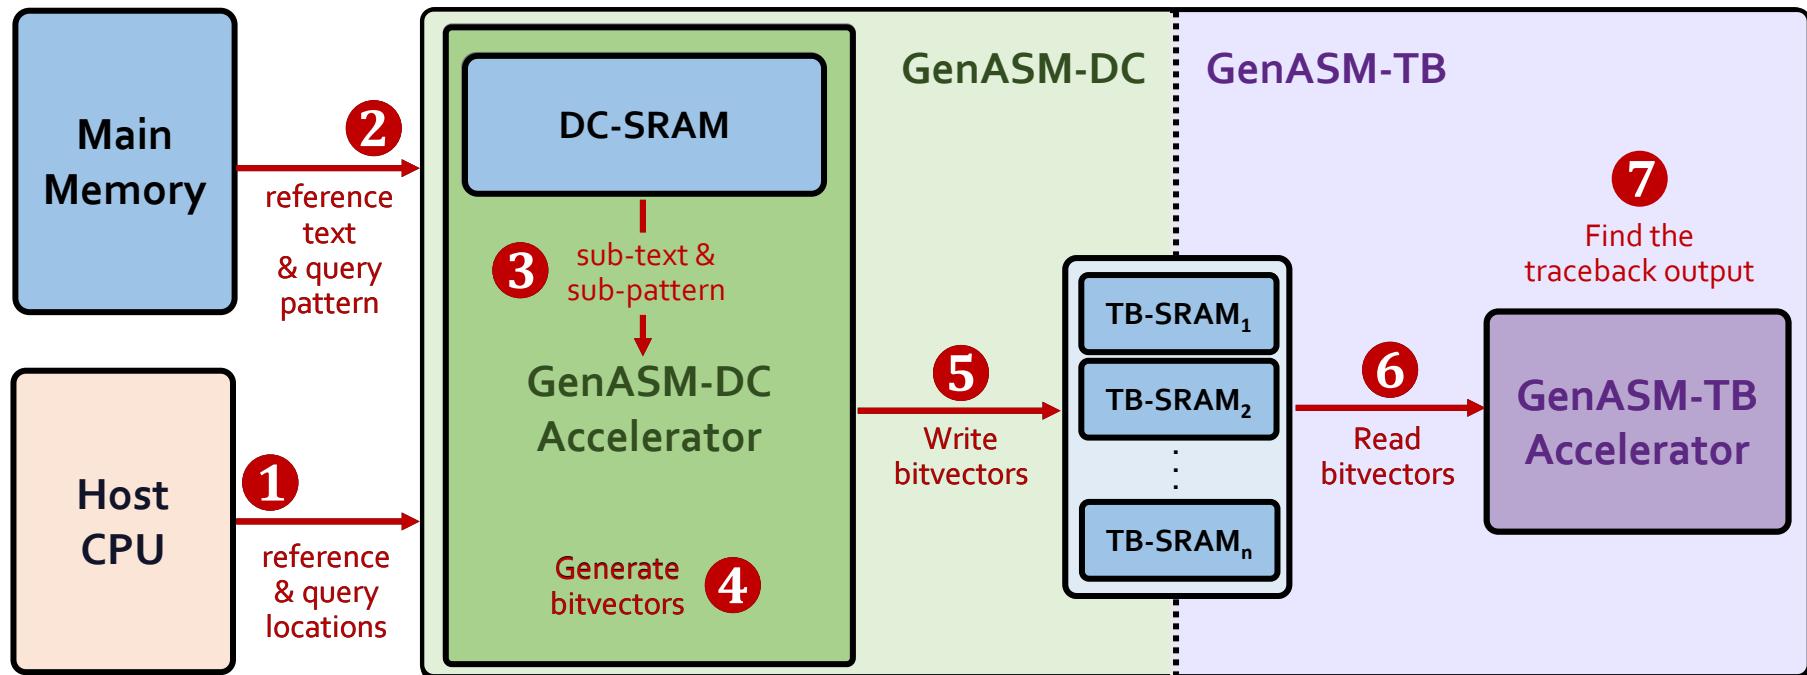
# GenASM Hardware Design



**GenASM-DC:**  
generates bitvectors  
and performs edit  
Distance Calculation

**GenASM-TB:**  
performs TraceBack  
and assembles the  
optimal alignment

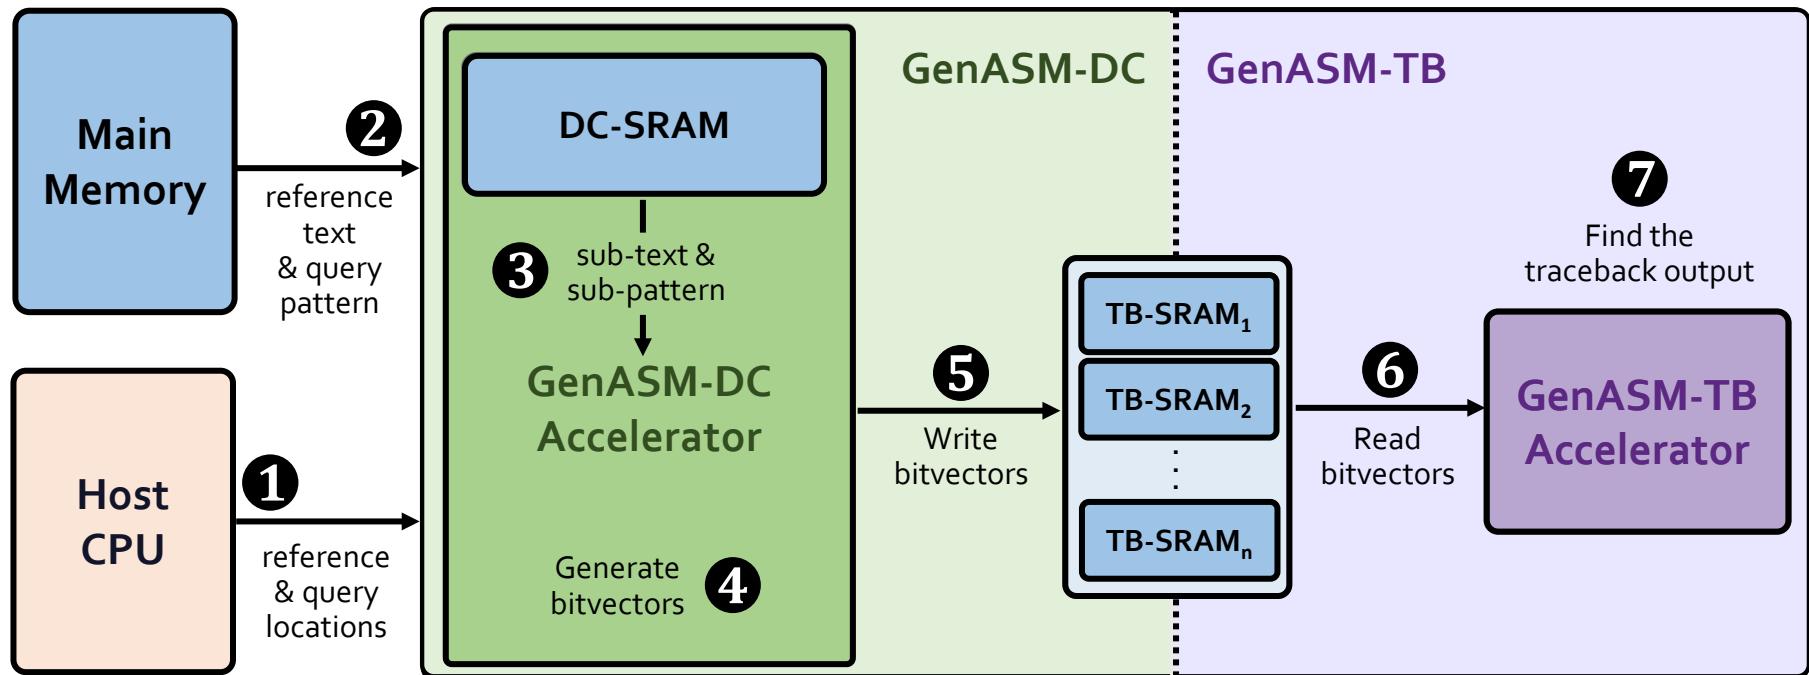
# GenASM Hardware Design



**GenASM-DC:**  
generates bitvectors  
and performs edit  
Distance Calculation

**GenASM-TB:**  
performs TraceBack  
and assembles the  
optimal alignment

# GenASM Hardware Design

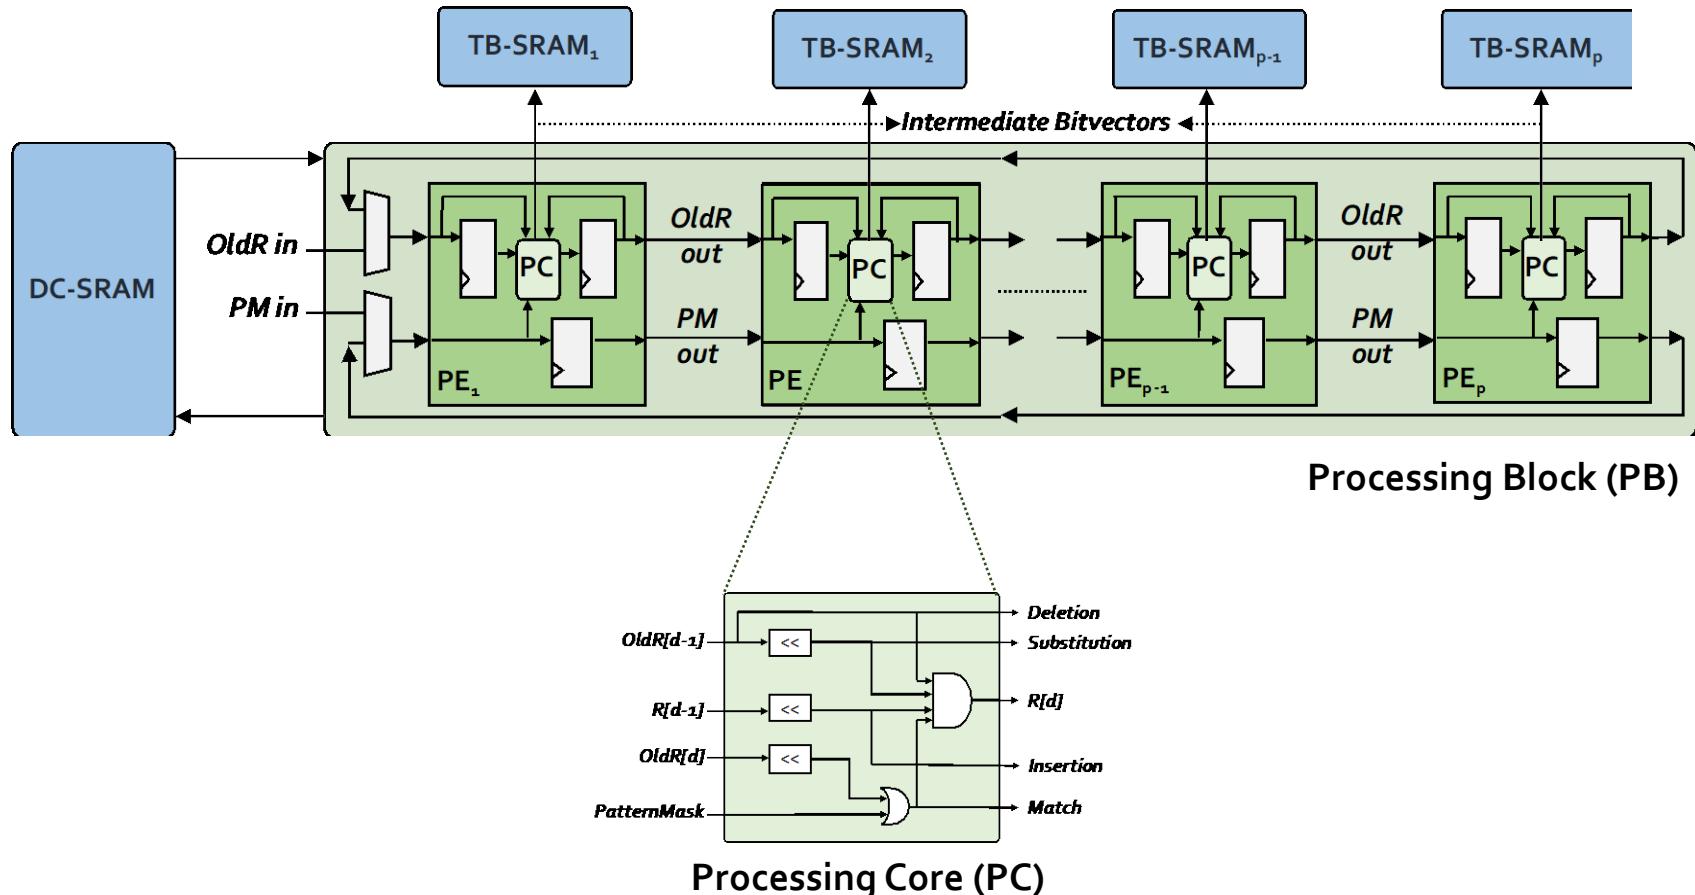


Our *specialized compute units* and *on-chip SRAMs* help us to:

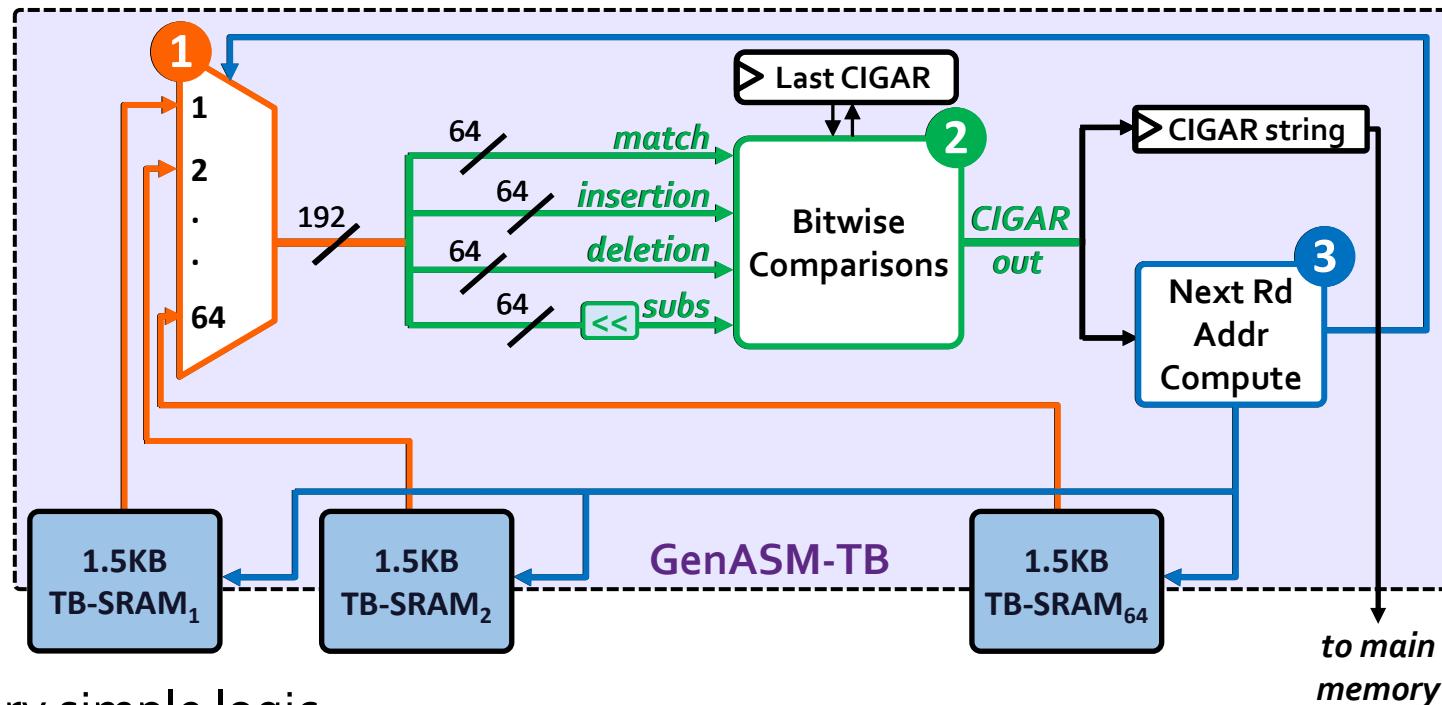
- Match **the rate of computation** with **memory capacity and bandwidth**
  - Achieve high performance and power efficiency
    - Scale linearly in performance with the number of parallel compute units that we add to the system

# GenASM-DC: Hardware Design

- Linear cyclic systolic array-based accelerator
  - Designed to maximize parallelism and minimize memory bandwidth and memory footprint



# GenASM-TB: Hardware Design



❑ Very simple logic:

- ① Reads the bitvectors from one of the TB-SRAMs using the computed address
- ② Performs the required bitwise comparisons to find the traceback output for the current position
- ③ Computes the next TB-SRAM address to read the new set of bitvectors

# Use Cases of GenASM

---

## (1) Read Alignment Step of Read Mapping

- Find the optimal alignment of how reads map to candidate reference regions

## (2) Pre-Alignment Filtering for Short Reads

- Quickly identify and filter out the unlikely candidate reference regions for each read

## (3) Edit Distance Calculation

- Measure the similarity or distance between two sequences
- We also discuss other possible use cases of GenASM in our paper:
  - Read-to-read overlap finding, hash-table based indexing, whole genome alignment, generic text search

# Evaluation Methodology

---

- We evaluate GenASM using:
  - Synthesized SystemVerilog models of the GenASM-DC and GenASM-TB accelerator datapaths
  - Detailed simulation-based performance modeling
  
- 16GB HMC-like 3D-stacked DRAM architecture
  - 32 vaults
  - 256GB/s of internal bandwidth, clock frequency of 1.25GHz
  - In order to achieve high parallelism and low power-consumption
  - Within each vault, the logic layer contains a GenASM-DC accelerator, its associated DC-SRAM, a GenASM-TB accelerator, and TB-SRAMs.

# Evaluation Methodology (cont'd.)

	SW Baselines	HW Baselines
Read Alignment	Minimap2 <sup>1</sup> BWA-MEM <sup>2</sup>	GACT (Darwin) <sup>3</sup> SillaX (GenAx) <sup>4</sup>
Pre-Alignment Filtering	–	Shouji <sup>5</sup>
Edit Distance Calculation	Edlib <sup>6</sup>	ASAP <sup>7</sup>

[1] H. Li. "Minimap2: Pairwise Alignment for Nucleotide Sequences." In *Bioinformatics*, 2018.

[2] H. Li. "Aligning sequence reads, clone sequences and assembly contigs with BWA-MEM." In *arXiv*, 2013.

[3] Y. Turakhia et al. "Darwin: A genomics co-processor provides up to 15,000 x acceleration on long read assembly." In *ASPLOS*, 2018.

[4] D. Fujiki et al. "GenAx: A genome sequencing accelerator." In *ISCA*, 2018.

[5] M. Alser. "Shouji: A fast and efficient pre-alignment filter for sequence alignment." In *Bioinformatics*, 2019.

[6] M. Šošić et al. "Edlib: A C/C++ library for fast, exact sequence alignment using edit distance." In *Bioinformatics*, 2017.

[7] S.S. Banerjee et al. "ASAP: Accelerated short-read alignment on programmable hardware." In *TC*, 2018.

# Evaluation Methodology (cont'd.)

---

- ❑ For Use Case 1: Read Alignment, we compare GenASM with:
  - Minimap2 and BWA-MEM (state-of-the-art **SW**)
    - Running on Intel® Xeon® Gold 6126 CPU (12-core) operating @2.60GHz with 64GB DDR4 memory
    - Using two simulated datasets:
      - Long ONT and PacBio reads: 10Kbp reads, 10-15% error rate
      - Short Illumina reads: 100-250bp reads, 5% error rate
  - GACT of Darwin and SillaX of GenAx (state-of-the-art **HW**)
    - Open-source RTL for GACT
    - Data reported by the original work for SillaX
    - GACT is best for long reads, SillaX is best for short reads

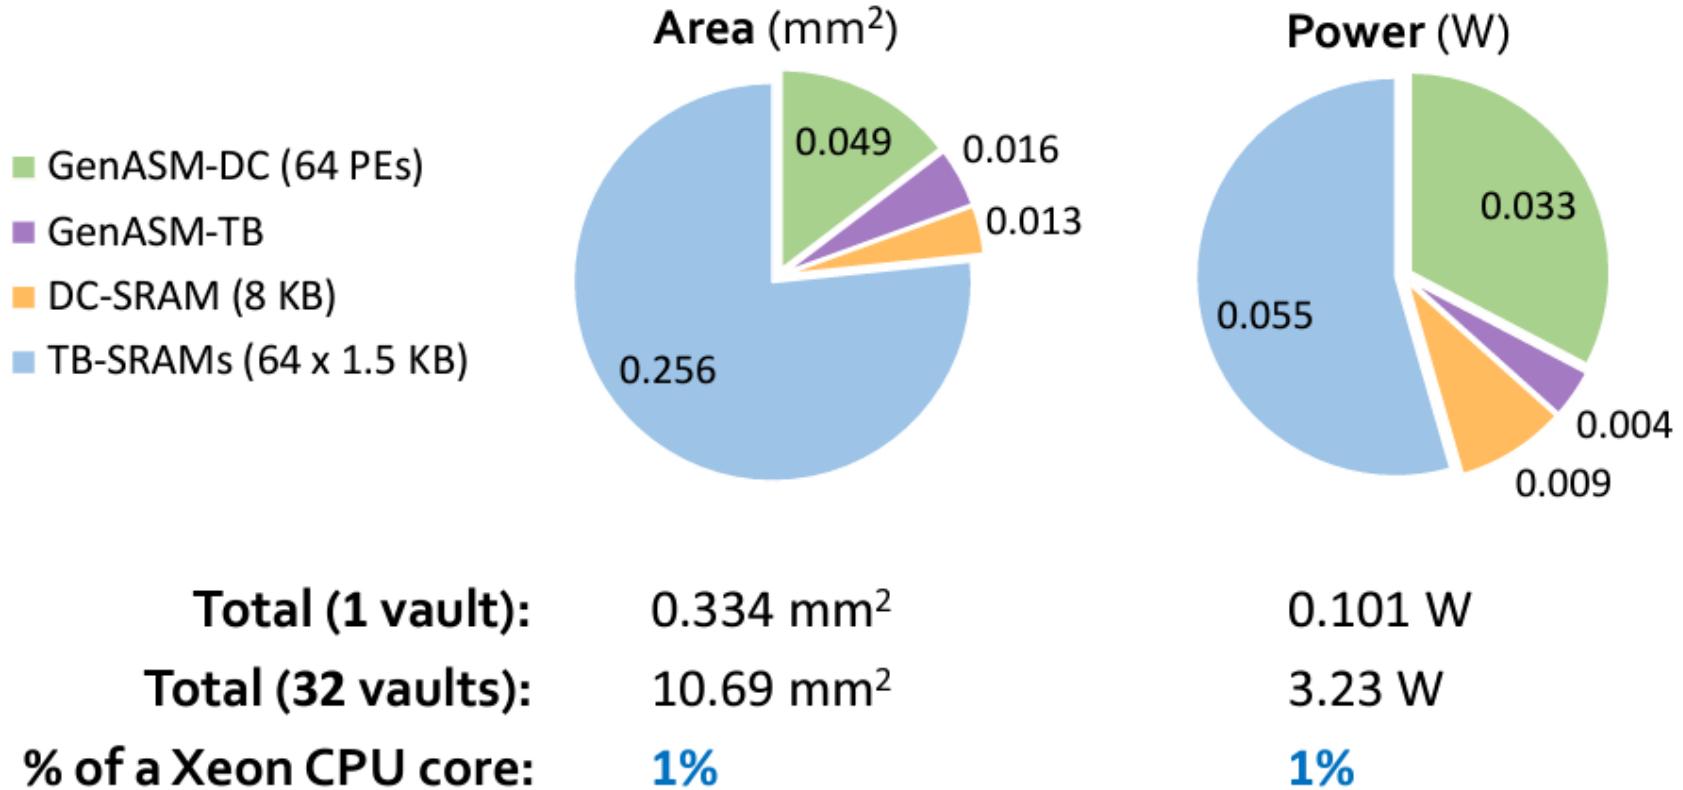
# Evaluation Methodology (cont'd.)

---

- ❑ For Use Case 2: Pre-Alignment Filtering, we compare GenASM with:
  - Shouji (state-of-the-art **HW** – FPGA-based filter)
    - Using two datasets provided as test cases:
      - 100bp reference-read pairs with an edit distance threshold of 5
      - 250bp reference-read pairs with an edit distance threshold of 15
  
- ❑ For Use Case 3: Edit Distance Calculation, we compare GenASM with:
  - Edlib (state-of-the-art **SW**)
    - Using two 100Kbp and 1Mbp sequences with similarity ranging between 60%-99%
  - ASAP (state-of-the-art **HW** – FPGA-based accelerator)
    - Using data reported by the original work

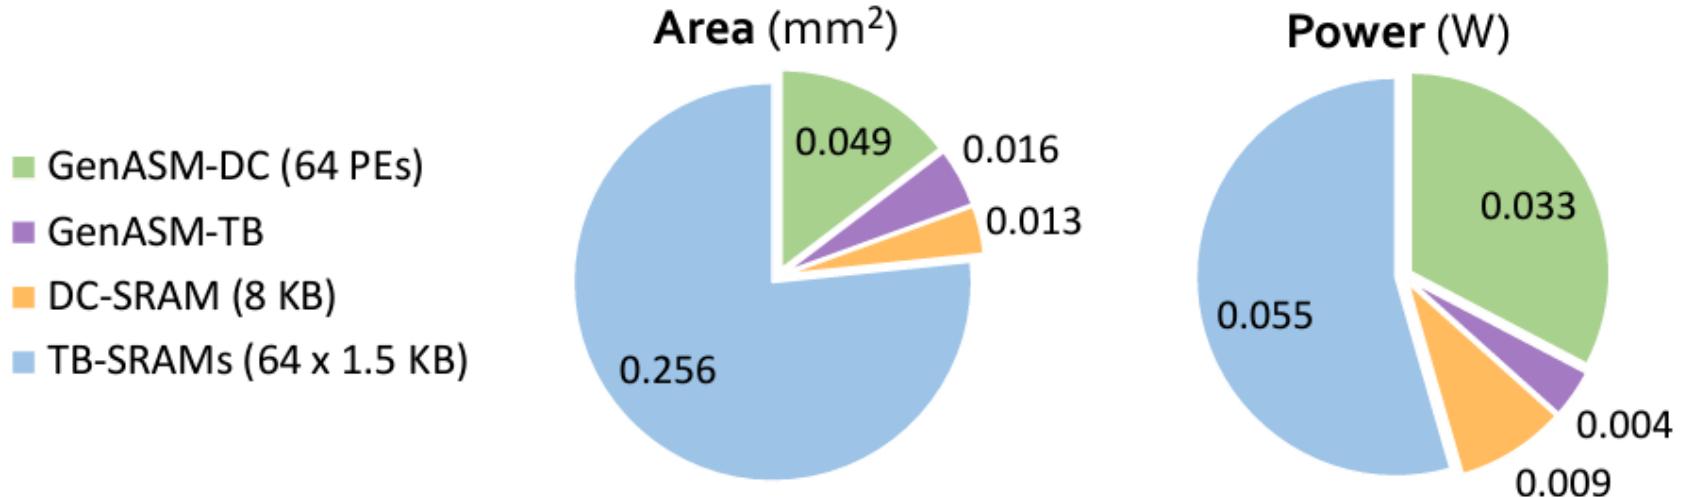
# Key Results – Area and Power

- Based on our **synthesis** of **GenASM-DC** and **GenASM-TB** accelerator datapaths using the Synopsys Design Compiler with a **28nm** process:
  - Both GenASM-DC and GenASM-TB operate **@ 1GHz**



# Key Results – Area and Power

- Based on our **synthesis** of **GenASM-DC** and **GenASM-TB** accelerator datapaths using the Synopsys Design Compiler with a **28nm** process:
  - Both GenASM-DC and GenASM-TB operate **@ 1GHz**



**GenASM has low area and power overheads**

# Key Results

## (1) Read Alignment

- ❑ **116×** speedup, **37×** less power than **Minimap2** (state-of-the-art **SW**)
- ❑ **111×** speedup, **33×** less power than **BWA-MEM** (state-of-the-art **SW**)
- ❑ **3.9×** better throughput, **2.7×** less power than **Darwin** (state-of-the-art **HW**)
- ❑ **1.9×** better throughput, **82%** less logic power than **GenAx** (state-of-the-art **HW**)

## (2) Pre-Alignment Filtering

- ❑ **3.7×** speedup, **1.7×** less power than **Shouji** (state-of-the-art **HW**), while significantly improving the accuracy of pre-alignment filtering

## (3) Edit Distance Calculation

- ❑ **22–12501×** speedup, **548–582×** less power than **Edlib** (state-of-the-art **SW**)
- ❑ **9.3–400×** speedup, **67×** less power than **ASAP** (state-of-the-art **HW**)

# Additional Details in the Paper

---

- Details of the **GenASM-DC and GenASM-TB algorithms**
- **Big-O analysis** of the algorithms
- Detailed explanation of **evaluated use cases**
- **Evaluation methodology details**  
(datasets, baselines, performance model)
- **Additional results** for the three evaluated use cases
- **Sources of improvements in GenASM**  
(algorithm-level, hardware-level, technology-level)
- Discussion of **four other potential use cases** of GenASM

# Summary of GenASM

---

- **Problem:**
  - Genome sequence analysis is bottlenecked by the **computational power** and **memory bandwidth limitations** of existing systems
  - This bottleneck is particularly an issue for **approximate string matching**
- **Key Contributions:**
  - GenASM: An approximate string matching (ASM) acceleration framework to accelerate **multiple steps of genome sequence analysis**
    - *First* to enhance and accelerate Bitap for ASM with genomic sequences
    - *Co-design* of our modified **scalable** and **memory-efficient** algorithms with **low-power** and **area-efficient** hardware accelerators
    - Evaluation of three different use cases: **read alignment**, **pre-alignment filtering**, **edit distance calculation**
- **Key Results:** GenASM is **significantly more efficient** for all the three use cases (in terms of **throughput** and **throughput per unit power**) than state-of-the-art **software** and **hardware** baselines

# GenASM [MICRO 2020]

---

Damla Senol Cali, Gurpreet S. Kalsi, Zulal Bingol, Can Firtina, Lavanya Subramanian, Jeremie S. Kim, Rachata Ausavarungnirun, Mohammed Alser, Juan Gomez-Luna, Amirali Boroumand, Anant Nori, Allison Scibisz, Sreenivas Subramoney, Can Alkan, Saugata Ghose, and Onur Mutlu,

"GenASM: A High-Performance, Low-Power Approximate String Matching Acceleration Framework for Genome Sequence Analysis"

*Proceedings of the 53rd International Symposium on Microarchitecture (MICRO), Virtual, October 2020.*

## GenASM: A High-Performance, Low-Power Approximate String Matching Acceleration Framework for Genome Sequence Analysis

Damla Senol Cali<sup>†✉</sup> Gurpreet S. Kalsi<sup>✉</sup> Zülal Bingöl<sup>▽</sup> Can Firtina<sup>◊</sup> Lavanya Subramanian<sup>‡</sup> Jeremie S. Kim<sup>◊†</sup>  
Rachata Ausavarungnirun<sup>○</sup> Mohammed Alser<sup>◊</sup> Juan Gomez-Luna<sup>◊</sup> Amirali Boroumand<sup>†</sup> Anant Nori<sup>✉</sup>  
Allison Scibisz<sup>†</sup> Sreenivas Subramoney<sup>✉</sup> Can Alkan<sup>▽</sup> Saugata Ghose<sup>\*†</sup> Onur Mutlu<sup>◊†▽</sup>

<sup>†</sup>Carnegie Mellon University    <sup>✉</sup>Processor Architecture Research Lab, Intel Labs    <sup>▽</sup>Bilkent University    <sup>◊</sup>ETH Zürich

<sup>‡</sup>Facebook    <sup>○</sup>King Mongkut's University of Technology North Bangkok    <sup>\*</sup>University of Illinois at Urbana-Champaign

# GenASM – GitHub Page

<https://github.com/CMU-SAFARI/GenASM>

The screenshot shows the GitHub repository page for 'CMU-SAFARI/GenASM'. The repository is public and has 1 branch and 0 tags. The master branch has 39 commits from 'damlasenolcali' (134ab9e on Apr 19). The repository description is: 'Source code for the software implementations of the GenASM algorithms proposed in our MICRO 2020 paper: Senol Cali et. al., "GenASM: A High-Performance, Low-Power Approximate String Matching Acceleration Framework for Genome Sequence Analysis" at https://people.inf.ethz.ch/omutlu/pub/GenASM-approximate-string-matching-framework-for-genome-analys...'. The repository has 26 stars, 5 watchers, and 6 forks. The README.md file is present.

Code

Issues Pull requests Actions Projects Wiki Security Insights Settings

master 1 branch 0 tags Go to file Add file Code About

damlasenolcali Update README.md 134ab9e on Apr 19 39 commits

LICENSE Initial commit 2 years ago

README.md Update README.md 2 months ago

genasm\_aligner.c Add files via upload 2 years ago

genasm\_aligner\_withDQ.c Add files via upload 15 months ago

genasm\_filter.c Add files via upload 2 years ago

README.md

## GenASM: Approximate String Matching (ASM) Acceleration Framework for Genome Sequence Analysis

GenASM is an approximate string matching (ASM) acceleration framework for genome sequence analysis. GenASM is a fast, efficient, and flexible framework for both short and long reads, which can be used to accelerate multiple steps of the genome sequence analysis pipeline. We base GenASM upon the Bitap algorithm. Bitap uses only fast and simple bitwise operations to perform approximate string matching. To our knowledge, GenASM is the first work that enhances and accelerates Bitap.

approximate-string-matching read-mapping  
hw-sw-co-design read-alignment  
bitap-algorithm pre-alignment-filtering  
genome-sequence-analysis

Readme GPL-3.0 license  
26 stars 5 watching 6 forks

# Research Contributions

Bottleneck analysis of genome assembly pipeline for long reads

*[Briefings in Bioinformatics, 2018]*

GenASM: Approximate string matching framework for genome sequence analysis

*[MICRO 2020]*

BitMAC: FPGA-based near-memory acceleration of bitvector-based sequence alignment

*[Ongoing]*

**SeGram:** Universal genomic mapping accelerator for both sequence-to-graph and sequence-to-sequence mapping

*[ISCA 2022]*

# Genome Sequence Analysis

---

- ***Sequence-to-sequence mapping (traditional read mapping):***
  - Critical step in genome sequence analysis (GSA)
  - Maps **reads** collected from an individual to a known **linear reference genome sequence**
  - Well studied with many available **tools** and accelerators
- Recent works replace the linear reference sequence with a **graph-based representation of the reference genome (genome graph)**
  - Captures the **genetic variations** and **diversity** across many individuals in a population
- ***Sequence-to-graph mapping*** results in notable quality improvements in GSA
  - More difficult computational problem
  - Much smaller number of practical software tools currently available
  - No prior hardware design for graph-based GSA

# Genome Graphs

---

Genome graphs:

- Combine the **linear reference genome** with the **known genetic variations** in the entire population as a graph-based data structure
- Enable us to move away from aligning with a single linear reference genome (**reference bias**) and **more accurately express the genetic diversity** in a population

**Sequence #1:** ACGTACGT

ACGTACGT

# Genome Graphs

---

Genome graphs:

- Combine the **linear reference genome** with the **known genetic variations** in the entire population as a graph-based data structure
- Enable us to move away from aligning with a single linear reference genome (**reference bias**) and **more accurately express the genetic diversity** in a population

**Sequence #1:** ACGTACGT

**Sequence #2:** ACGGACGT

ACGTACGT

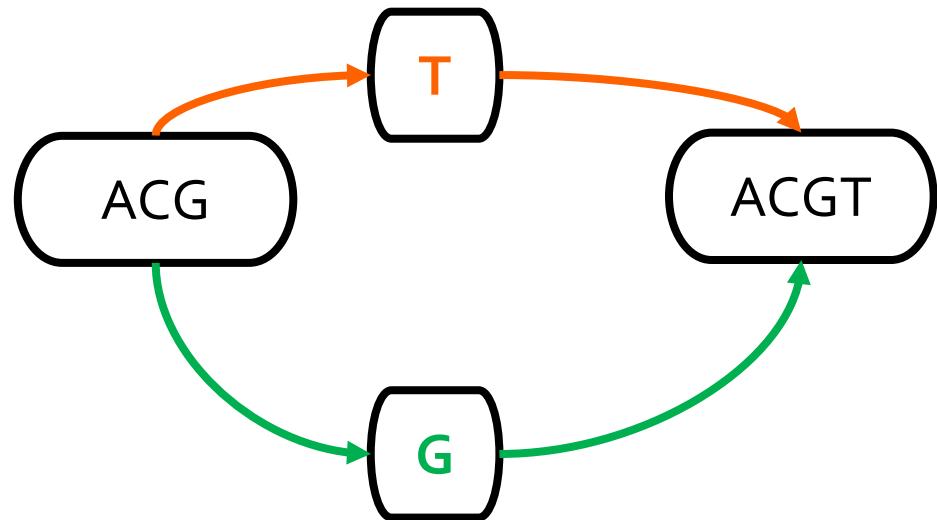
# Genome Graphs

Genome graphs:

- Combine the **linear reference genome** with the **known genetic variations** in the entire population as a graph-based data structure
- Enable us to move away from aligning with a single linear reference genome (**reference bias**) and **more accurately express the genetic diversity** in a population

**Sequence #1:** ACGTACGT

**Sequence #2:** ACGGACGT



# Genome Graphs

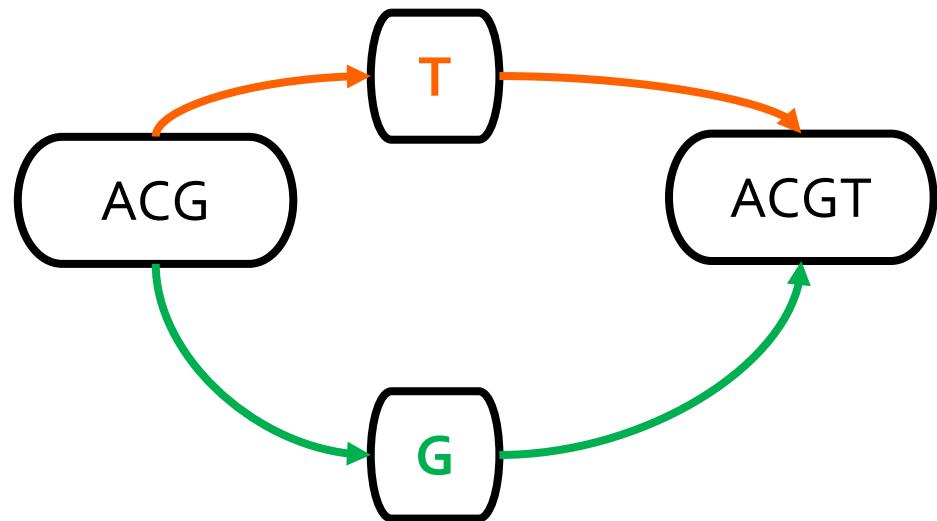
Genome graphs:

- Combine the **linear reference genome** with the **known genetic variations** in the entire population as a graph-based data structure
- Enable us to move away from aligning with a single linear reference genome (**reference bias**) and **more accurately express the genetic diversity** in a population

**Sequence #1:** ACGTACGT

**Sequence #2:** ACGGACGT

**Sequence #3:** ACGTTACGT



# Genome Graphs

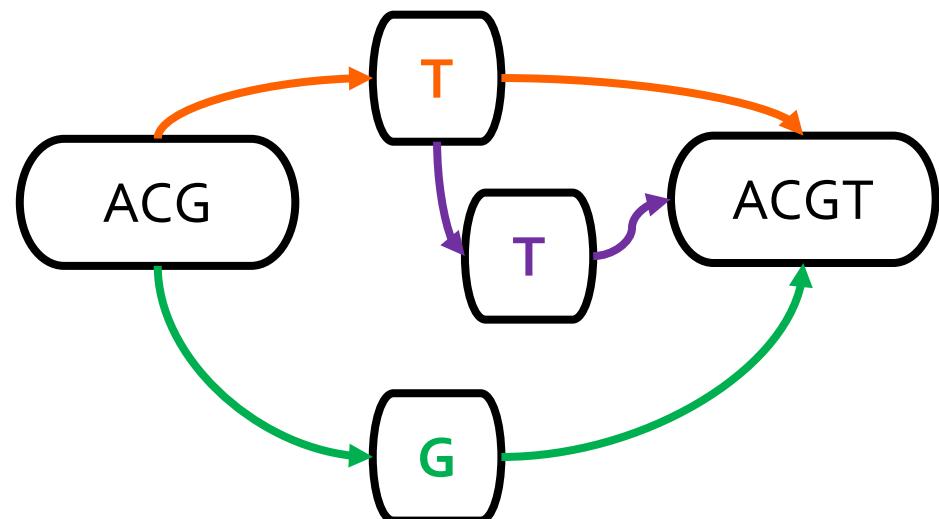
Genome graphs:

- Combine the **linear reference genome** with the **known genetic variations** in the entire population as a graph-based data structure
- Enable us to move away from aligning with a single linear reference genome (**reference bias**) and **more accurately express the genetic diversity** in a population

**Sequence #1:** ACGTACGT

**Sequence #2:** ACGGACGT

**Sequence #3:** ACGTTACGT



# Genome Graphs

Genome graphs:

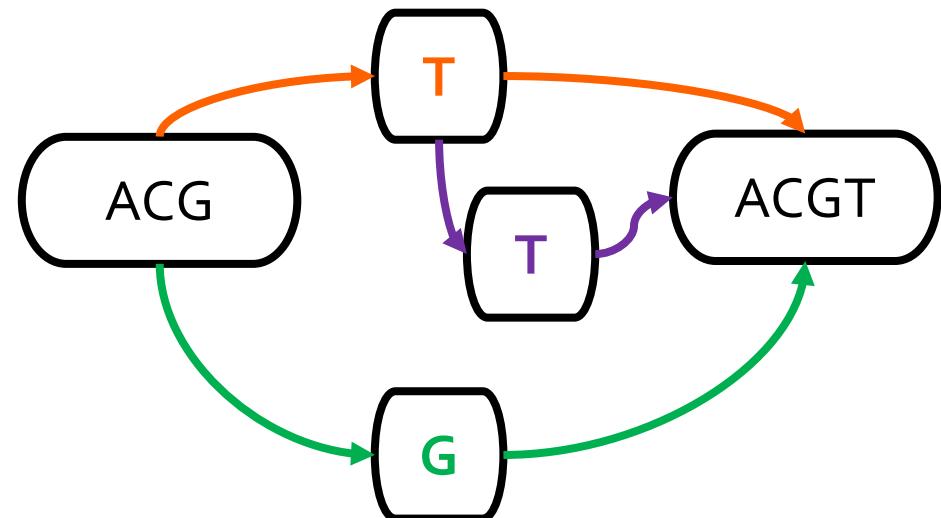
- Combine the **linear reference genome** with the **known genetic variations** in the entire population as a graph-based data structure
- Enable us to move away from aligning with a single linear reference genome (**reference bias**) and **more accurately express the genetic diversity** in a population

**Sequence #1:** ACGTACGT

**Sequence #2:** ACGGACGT

**Sequence #3:** ACGTTACGT

**Sequence #4:** ACGACGT



# Genome Graphs

Genome graphs:

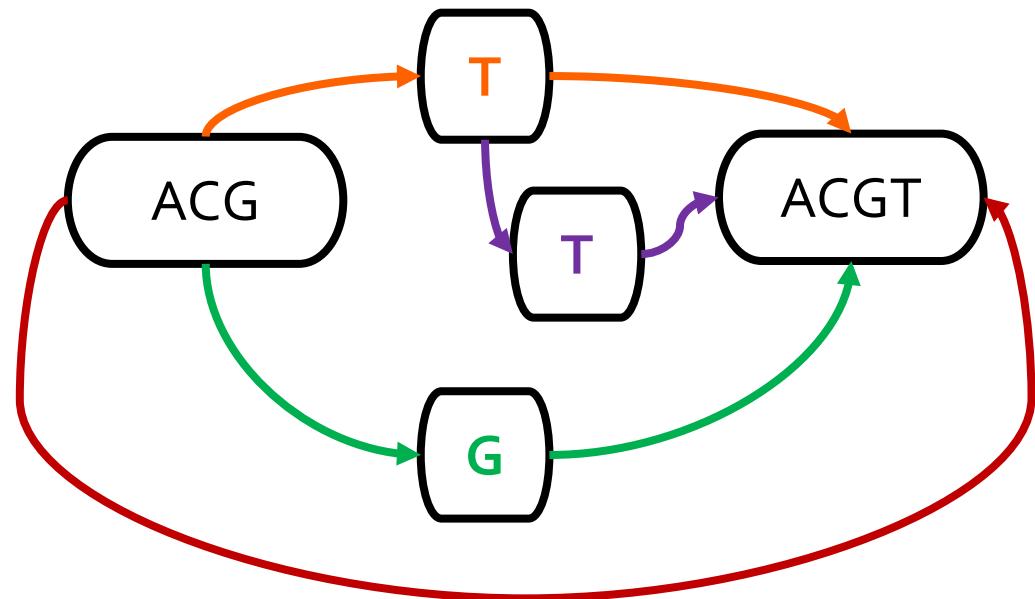
- Combine the **linear reference genome** with the **known genetic variations** in the entire population as a graph-based data structure
- Enable us to move away from aligning with a single linear reference genome (**reference bias**) and **more accurately express the genetic diversity** in a population

**Sequence #1:** ACGTACGT

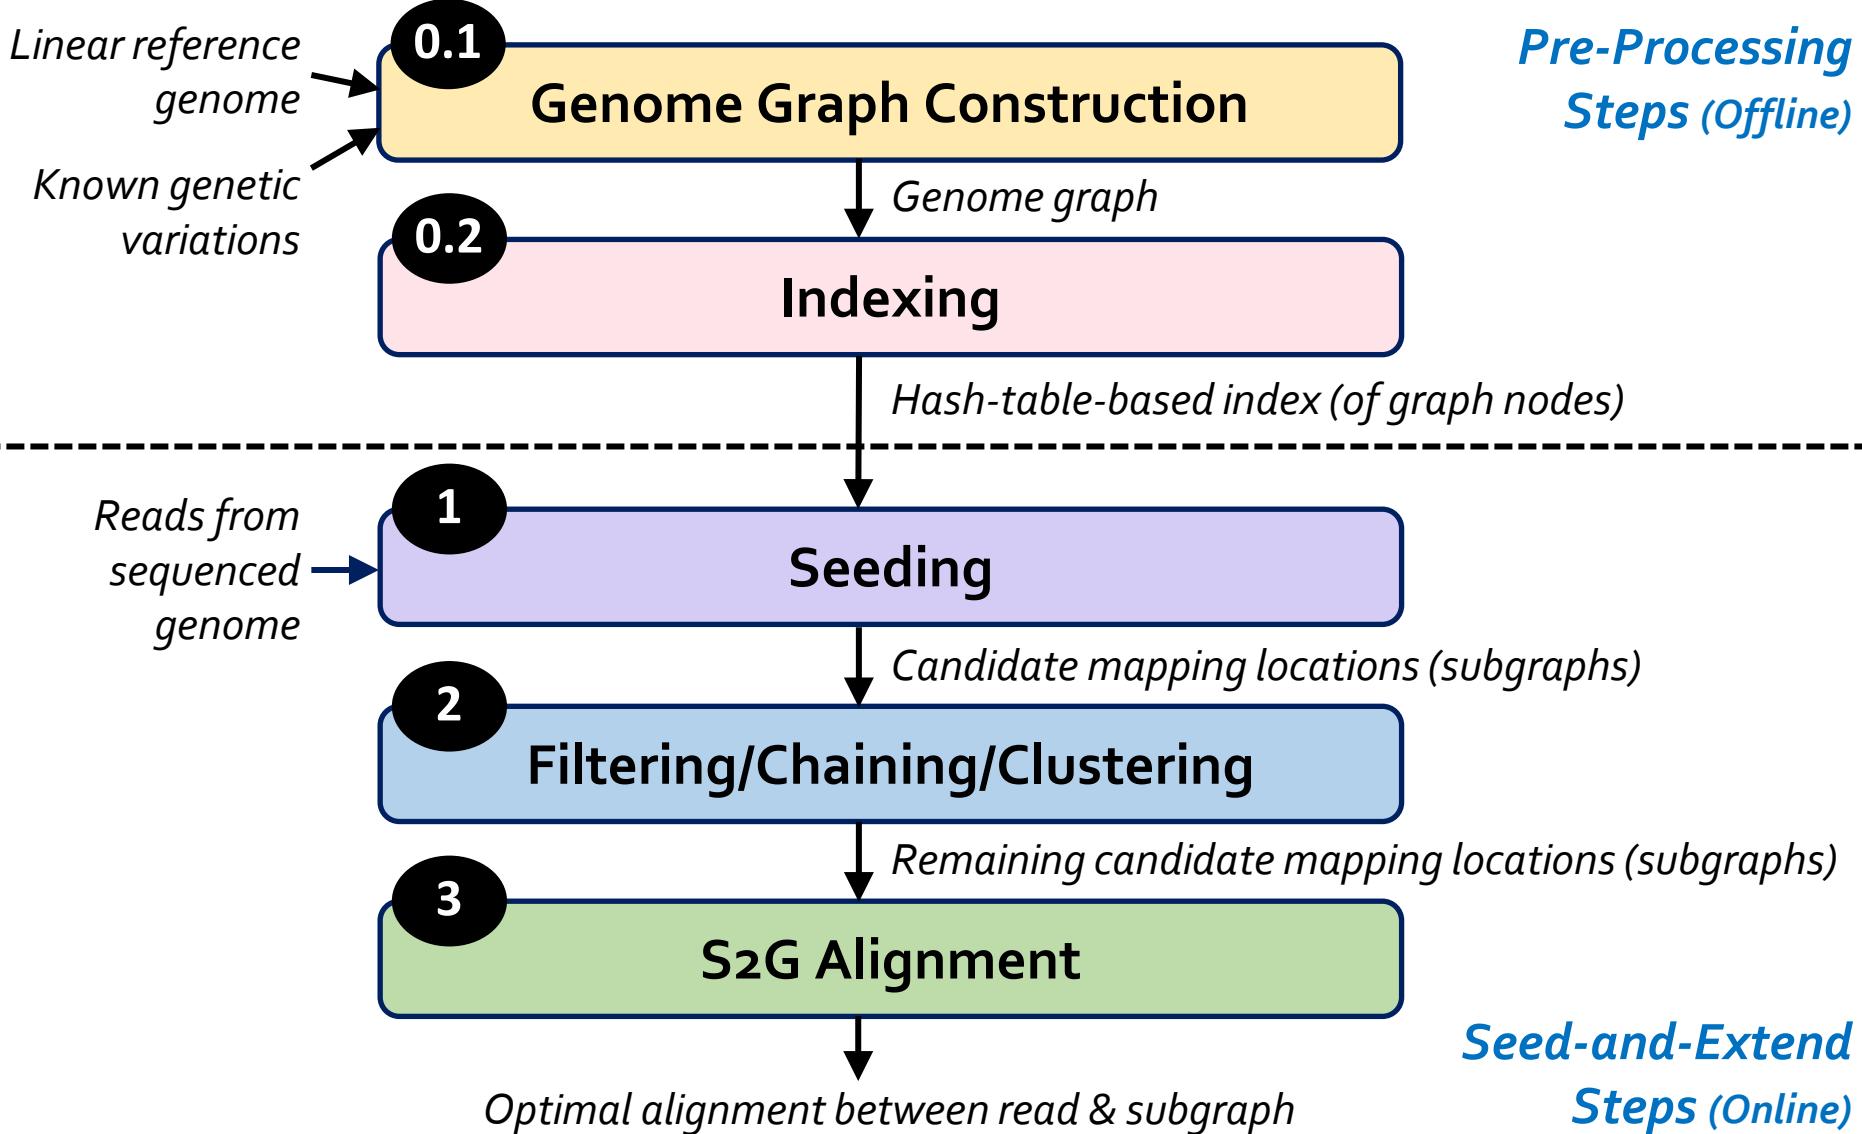
**Sequence #2:** ACGGACGT

**Sequence #3:** ACGTTACGT

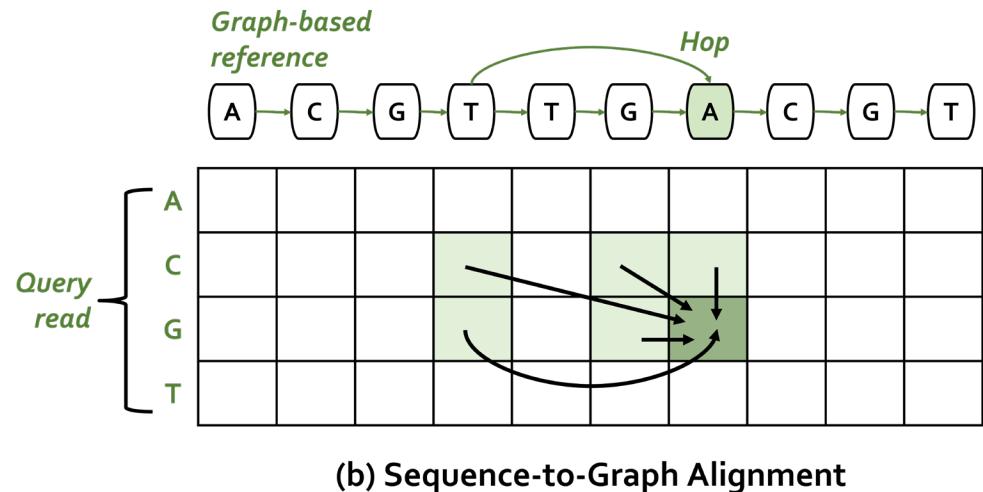
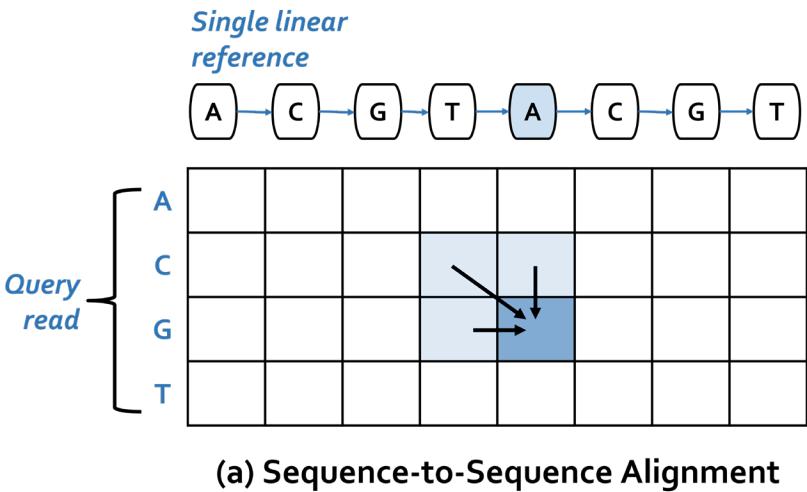
**Sequence #4:** ACGACGT



# Sequence-to-Graph Mapping Pipeline



# S<sub>2</sub>S vs. S<sub>2</sub>G Alignment



In contrast to S<sub>2</sub>S alignment,  
S<sub>2</sub>G alignment must incorporate **non-neighboring characters**  
as well whenever there is an edge (i.e., **hop**)  
from the non-neighboring character to the current character

# Analysis of State-of-the-Art Tools

**Observation 1:** Alignment Step is the Bottleneck

SW

**Observation 2:** Alignment Suffers from High Cache Miss Rates

**Observation 3:** Seeding Suffers from the DRAM Latency

Bottleneck

**Observation 4:** Baseline Tools Scale Sublinearly

**Observation 5:** Existing S2S Mapping Accelerators are  
Unsuitable for the S2G Mapping Problem

HW

**Observation 6:** Existing Graph Accelerators are Unable to  
Handle S2G Alignment

# SeGram: Universal Genomic Mapping Accelerator

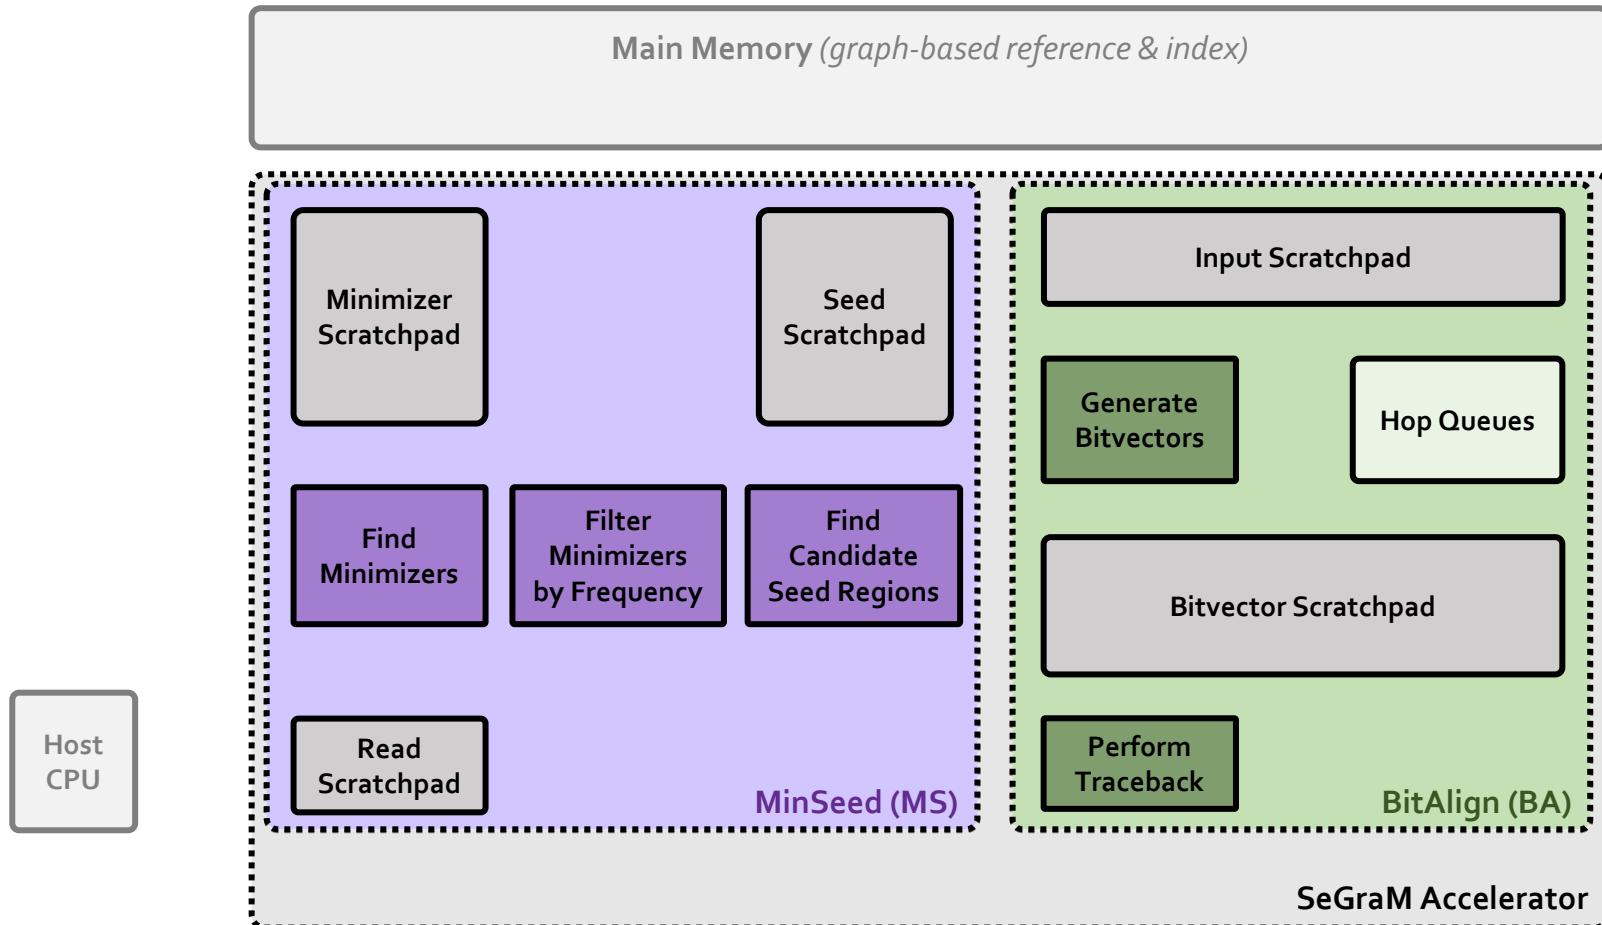
## Our Goal:

**Specialized, high-performance, scalable, and low-cost**  
algorithm/hardware co-design that alleviates bottlenecks in  
*both the seeding and alignment steps* of  
sequence-to-graph mapping

**SeGram:** *First universal genomic mapping accelerator* that can support both sequence-to-graph and sequence-to-sequence mapping, for both short and long reads

- First algorithm/hardware co-design*** for sequence-to-graph mapping
- We base SeGram upon a **minimizer-based seeding algorithm** and a **novel bitvector-based alignment algorithm** SW
- We co-design both algorithms with **high-performance, scalable, and efficient hardware accelerators** HW

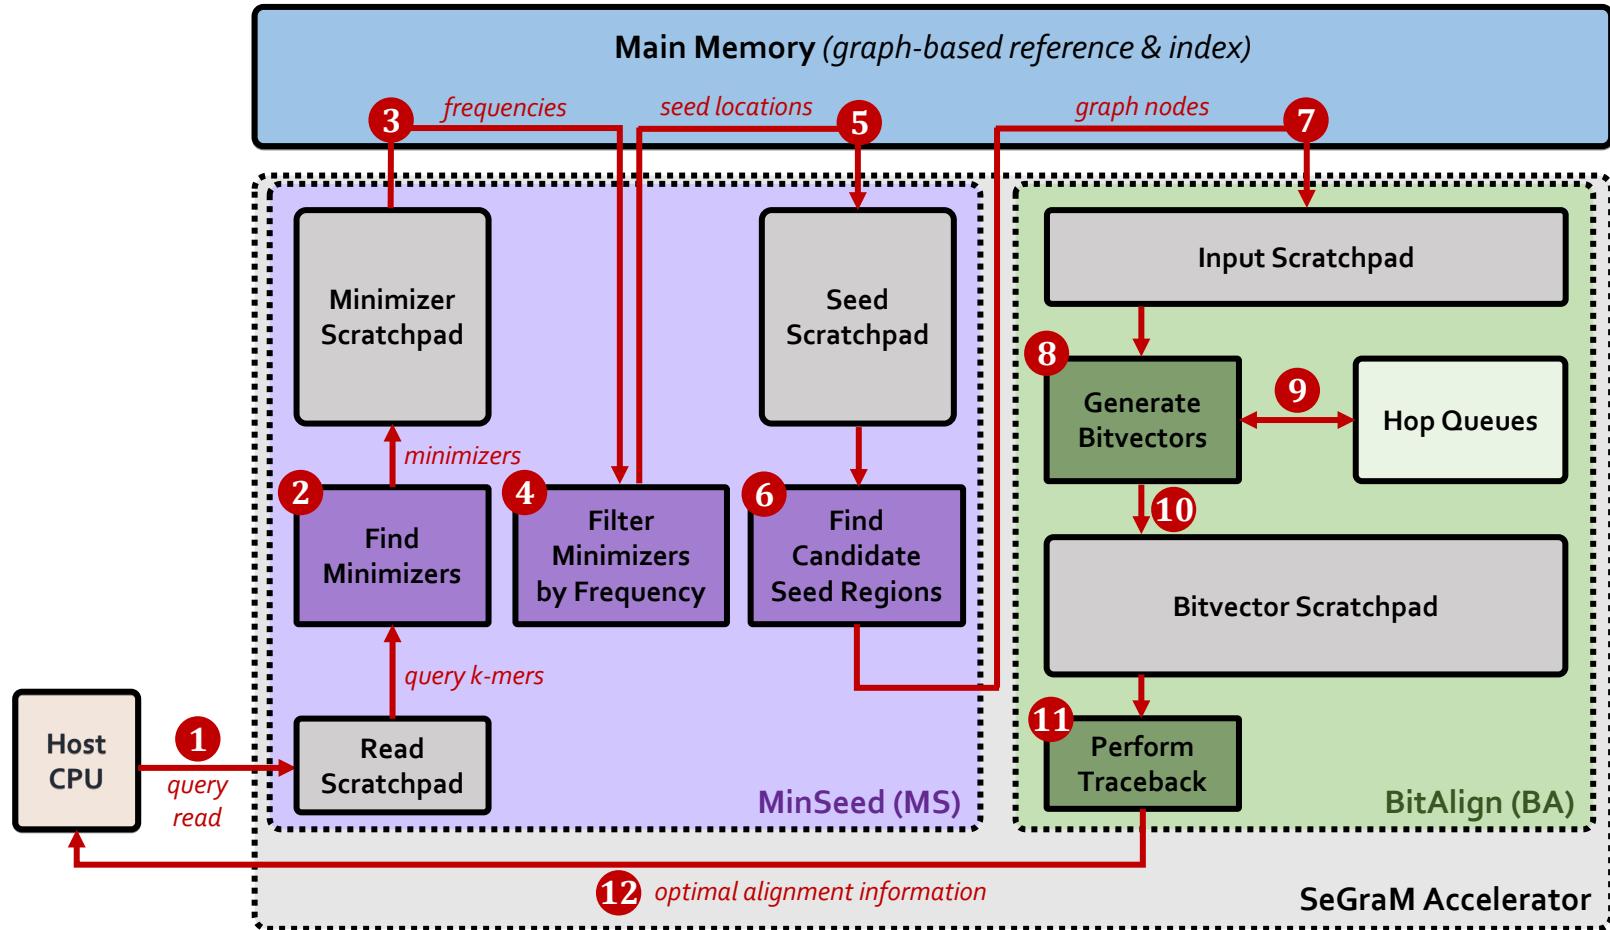
# SeGraM Hardware Design



**MinSeed:** first hardware  
accelerator for  
**Minimizer-based Seeding**

**BitAlign:** first hardware  
accelerator for (**Bitvector-based**)  
sequence-to-graph **Alignment**

# SeGram Hardware Design

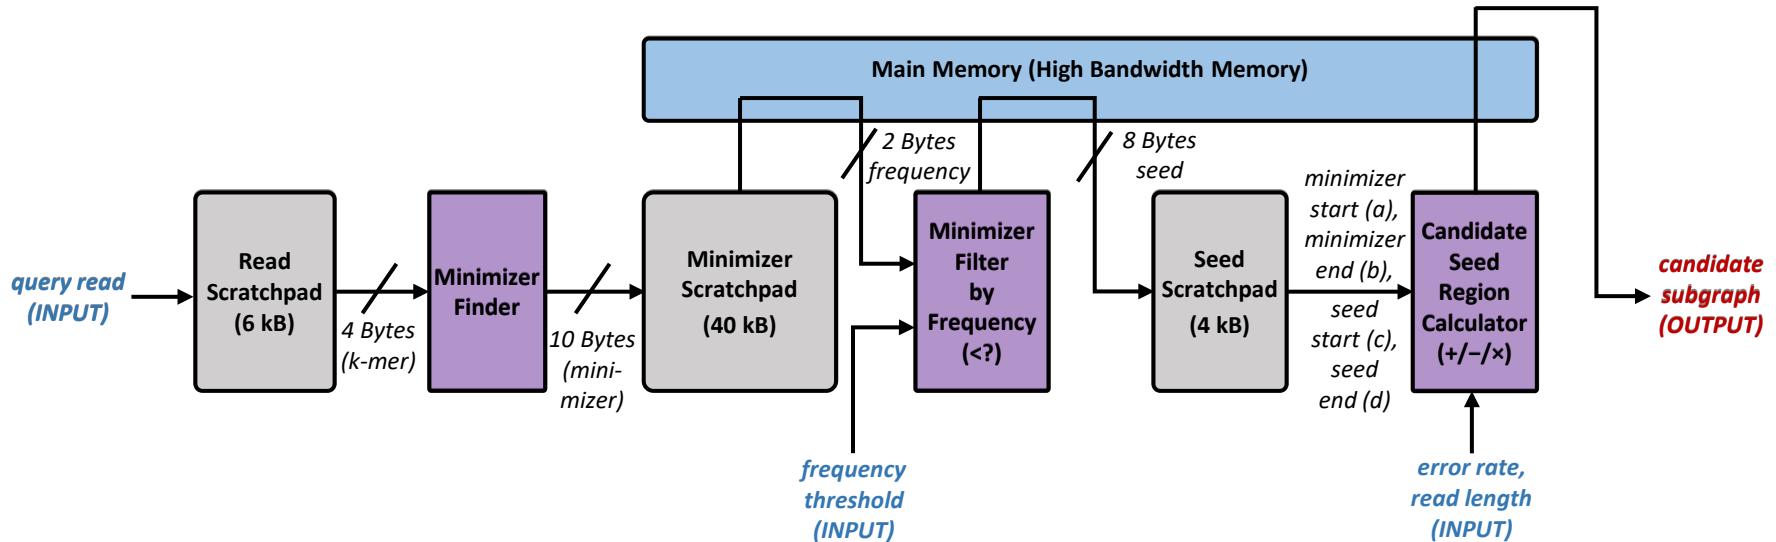


**MinSeed:** first hardware accelerator for Minimizer-based Seeding

**BitAlign:** first hardware accelerator for (Bitvector-based) sequence-to-graph Alignment

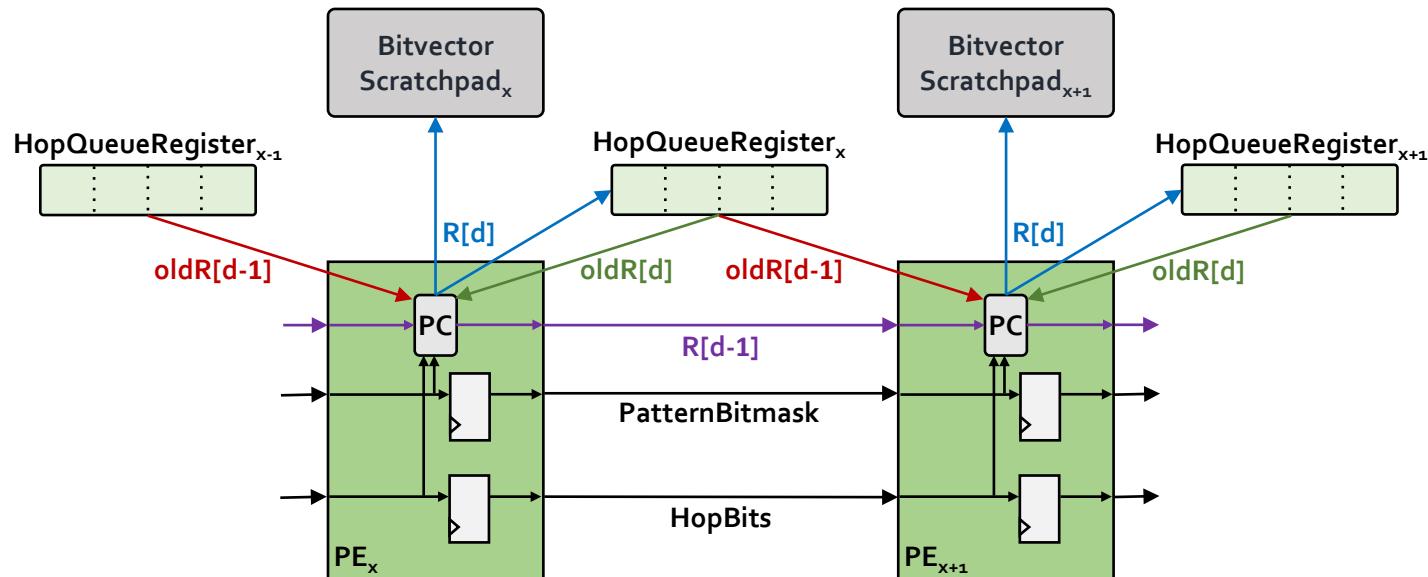
# MinSeed HW

- MinSeed accelerator consists of **three computation modules**, three scratchpads, and **the memory interface**
  - **Computation modules** are implemented with simple logic
  - For all three scratchpads, we employ a double buffering technique to hide the latency of the MinSeed accelerator
  - We couple MinSeed with **High-Bandwidth Memory (HBM)** to enable low-latency and highly-parallel memory access

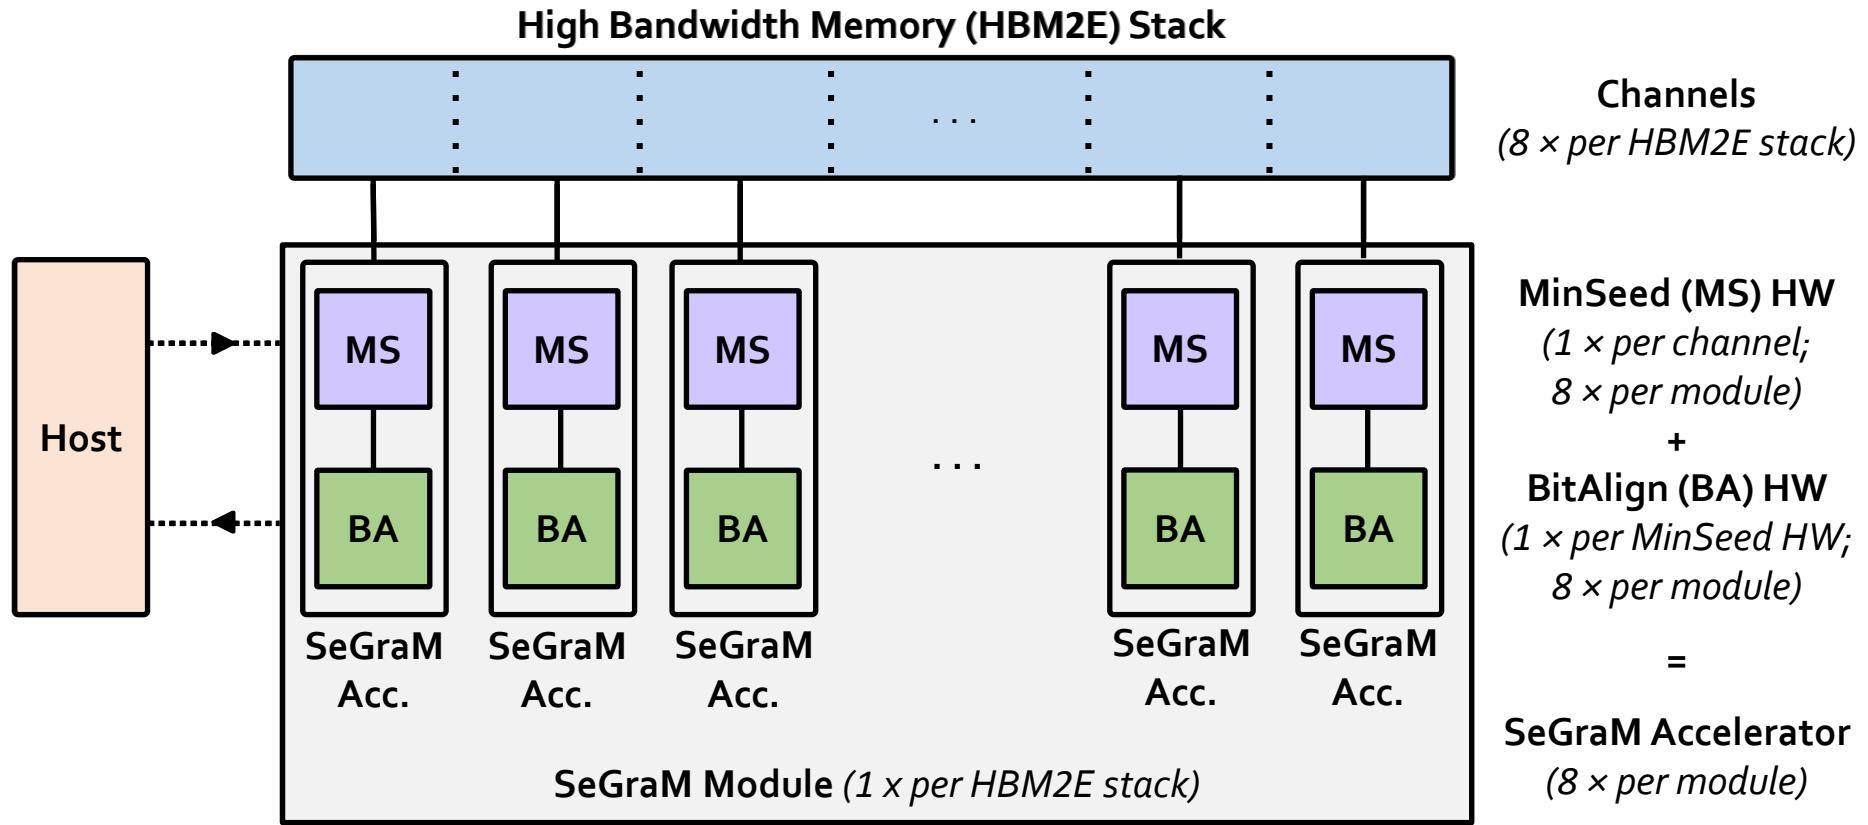


# BitAlign HW

- ❑ Linear cyclic systolic array-based accelerator
- ❑ Incorporates *hop queue registers* to feed the bitvectors of non-neighboring characters/nodes (i.e., *hops*)
- ❑ We implement the hop information between nodes of the graph as an adjacency matrix called **HopBits**
- ❑ Based on empirical analysis, we select **12** as the *hop limit*



# Overall System Design of SeGraM



# Use Cases of SeGraM

---

## (1) End-to-End Sequence-to-Graph Mapping

- The whole SeGraM design (MinSeed + BitAlign) should be employed
- We can use SeGraM to perform mapping with both short and long reads

## (2) Sequence-to-Graph Alignment

- BitAlign can be used as a standalone sequence-to-graph aligner without the need of an initial seeding tool/accelerator (e.g., MinSeed)
- BitAlign is orthogonal to and can be coupled with any seeding (or filtering) tool/accelerator

## (3) Sequence-to-Sequence Alignment

- BitAlign can also be used for sequence-to-sequence alignment, as it is a special and simpler variant of sequence-to-graph alignment

## (4) Seeding

- MinSeed can be used as a standalone seeding accelerator for both graph-based mapping and traditional linear mapping
- MinSeed is orthogonal to and can be coupled with any alignment tool/accelerator

# Evaluation Methodology

---

## ❑ Performance, Area and Power Analysis:

- **Synthesized** SystemVerilog models of the MinSeed and BitAlign accelerator datapaths
- **Simulation- and spreadsheet-based** performance modeling

## ❑ Baseline Comparison Points:

- **GraphAligner, vg, and HGA** for sequence-to-graph mapping
- **PaSGAL** for sequence-to-graph alignment
- **Darwin, GenAx, and GenASM** for sequence-to-sequence alignment

## ❑ Datasets:

- **Graph-based reference:** GRCh38 + 7 VCF files for HG001-007
- **Simulated datasets** for both short and long reads

# Key Results – Area & Power

- Based on our **synthesis** of **MinSeed** and **BitAlign** accelerator datapaths using the Synopsys Design Compiler with a **28nm** process (**@ 1GHz**):

Component	Area (mm <sup>2</sup> )	Power (mW)
MinSeed – Logic	0.017	10.8
Read Scratchpad (6 kB)	0.012	7.9
Minimizer Scratchpad (40 kB)	0.055	22.7
Seed Scratchpad (4 kB)	0.008	6.4
BitAlign – Edit Distance Calculation Logic with Hop Queue Registers (64 PEs)	0.393	378.0
BitAlign – Traceback Logic	0.020	2.7
Input Scratchpad (24 kB)	0.033	13.3
Bitvector Scratchpads (128 kB)	0.329	316.2
<b>Total – 1 SeGram Accelerator</b>	<b>0.867</b>	<b>758.0 (0.8 W)</b>
<b>Total – 32 SeGram Accelerators</b>	<b>27.744</b>	<b>24256.0 (24.3 W)</b>
HBM2E (4 stacks)	--	3.8 W

# Key Results

## (1) Sequence-to-Graph (S<sub>2</sub>G) Mapping

- ❑ **5.9×/106×** speedup, **4.1×/3.0×** less power than **GraphAligner** for long and short reads, respectively (state-of-the-art **SW**)
- ❑ **3.9×/742×** speedup, **4.4×/3.2×** less power than **vg** for long and short reads, respectively (state-of-the-art **SW**)

## (2) Sequence-to-Graph (S<sub>2</sub>G) Alignment

- ❑ **41×–539×** speedup over **PaSGAL** with AVX-512 support (state-of-the-art **SW**)

## (3) Sequence-to-Sequence (S<sub>2</sub>S) Alignment

- ❑ **1.2×/4.8×** higher throughput than **GenASM** and **GACT** of **Darwin** for long reads (state-of-the-art **HW**)
- ❑ **1.3×/2.4×** higher throughput than **GenASM** and **SillaX** of **GenAX** for short reads (state-of-the-art **HW**)

# Additional Details in the Paper

---

- Details of the **pre-processing steps of SeGram**
- Details of the **MinSeed and BitAlign algorithms**
- **Bottleneck analysis** of the existing tools
- **Evaluation methodology details**  
(datasets, baselines, performance model)
- **Additional results** for the three evaluated use cases
- **Sources of improvements in SeGram**
- **Comparison of GenASM and SeGram**

# Summary of SeGram

---

## Problem:

- Sequence-to-sequence (S2S) mapping causes **reference bias**
- Recent works replace the linear reference sequence with a graph-based representation of the reference genome
- Sequence-to-graph (S2G) mapping is a **more difficult computational problem**, with a much smaller number of practical software tools

## Key Contributions:

- **SeGram**: *Universal algorithm/hardware co-designed genomic mapping accelerator that supports both S2G and S2S mapping*
  - **MinSeed**: *First minimizer-based seeding accelerator*
  - **BitAlign**: *First sequence-to-graph alignment accelerator based upon our new bitvector-based, highly-parallel algorithm*

## Key Results:

- **SeGram** provides **greatly higher throughput and lower power consumption** compared to state-of-the-art SW tools for S2G mapping
- **BitAlign significantly outperforms** a state-of-the-art S2G alignment tool and three state-of-the-art HW solutions for S2S alignment

# SeGram [ISCA 2022]

---

Damla Senol Cali, Konstantinos Kanellopoulos, Joel Lindegger, Zulal Bingol, Gurpreet S. Kalsi, Ziyi Zuo, Can Firtina, Meryem Banu Cavlak, Jeremie S. Kim, Nika Mansouri Ghiasi, Gagandeep Singh, Juan Gomez-Luna, Nour Almadhoun Alserr, Mohammed Alser, Sreenivas Subramoney, Can Alkan, Saugata Ghose, and Onur Mutlu

"SeGram: A Universal Hardware Accelerator for Genomic Sequence-to-Graph and Sequence-to-Sequence Mapping"

*Proceedings of the 49th International Symposium on Computer Architecture (ISCA),*  
New York City, NY, June 2022.

## **SeGram: A Universal Hardware Accelerator for Genomic Sequence-to-Graph and Sequence-to-Sequence Mapping**

Damla Senol Cali<sup>1,2</sup> Konstantinos Kanellopoulos<sup>2</sup> Joël Lindegger<sup>2</sup> Zülal Bingöl<sup>3</sup>  
Gurpreet S. Kalsi<sup>4</sup> Ziyi Zuo<sup>5</sup> Can Firtina<sup>2</sup> Meryem Banu Cavlak<sup>2</sup> Jeremie Kim<sup>2</sup>  
Nika Mansouri Ghiasi<sup>2</sup> Gagandeep Singh<sup>2</sup> Juan Gómez-Luna<sup>2</sup> Nour Almadhoun Alserr<sup>2</sup>  
Mohammed Alser<sup>2</sup> Sreenivas Subramoney<sup>4</sup> Can Alkan<sup>3</sup> Saugata Ghose<sup>6</sup> Onur Mutlu<sup>2</sup>

<sup>1</sup>Bionano Genomics   <sup>2</sup>ETH Zürich   <sup>3</sup>Bilkent University   <sup>4</sup>Intel Labs

<sup>5</sup>Carnegie Mellon University   <sup>6</sup>University of Illinois Urbana-Champaign

# SeGram – GitHub Page

<https://github.com/CMU-SAFARI/SeGram>

The screenshot shows the GitHub repository page for 'CMU-SAFARI / SeGram'. The repository is public, as indicated by the 'Public' badge. The main navigation bar includes links for Code, Issues, Pull requests, Actions, Projects, Wiki, Security, Insights, and Settings. The 'Code' tab is selected.

Key statistics displayed are: 1 branch, 0 tags, and 6 commits. The commit history shows:

- damlasenolcali Update README.md (0837f80, 2 days ago)
- LICENSE Initial commit (2 months ago)
- README.md Update README.md (2 days ago)

A large callout box highlights the README.md file content:

**SeGram (Software implementations and datasets will be available soon!)**

The detailed description of the software states:

SeGram is a universal genomic mapping accelerator that supports both sequence-to-graph mapping and sequence-to-sequence mapping, for both short and long reads. SeGram consists of two main components: (1) MinSeed, the first minimizer-based seeding accelerator, which finds the candidate mapping locations (i.e., subgraphs) in a given genome graph; and (2) BitAlign, the first bitvector-based sequence-to-graph alignment accelerator, which performs alignment between a given read and the subgraph identified by MinSeed. MinSeed is built upon a memory-efficient minimizer-based seeding algorithm, and BitAlign is built upon our novel bitvector-based, highly-parallel sequence-to-graph alignment algorithm.

The repository page also includes sections for About, Releases, and various metadata like Readme, MIT license, and star counts.

# Accelerating Genome Sequence Analysis via Efficient Hardware/Algorithm Co-Design

**Damlasenol Cali, Ph.D.**

<https://damlasenolcali.github.io/>

[damlasenolcali@gmail.com](mailto:damlasenolcali@gmail.com)

Staff Software Engineer, Hardware Acceleration  
Bionano Genomics

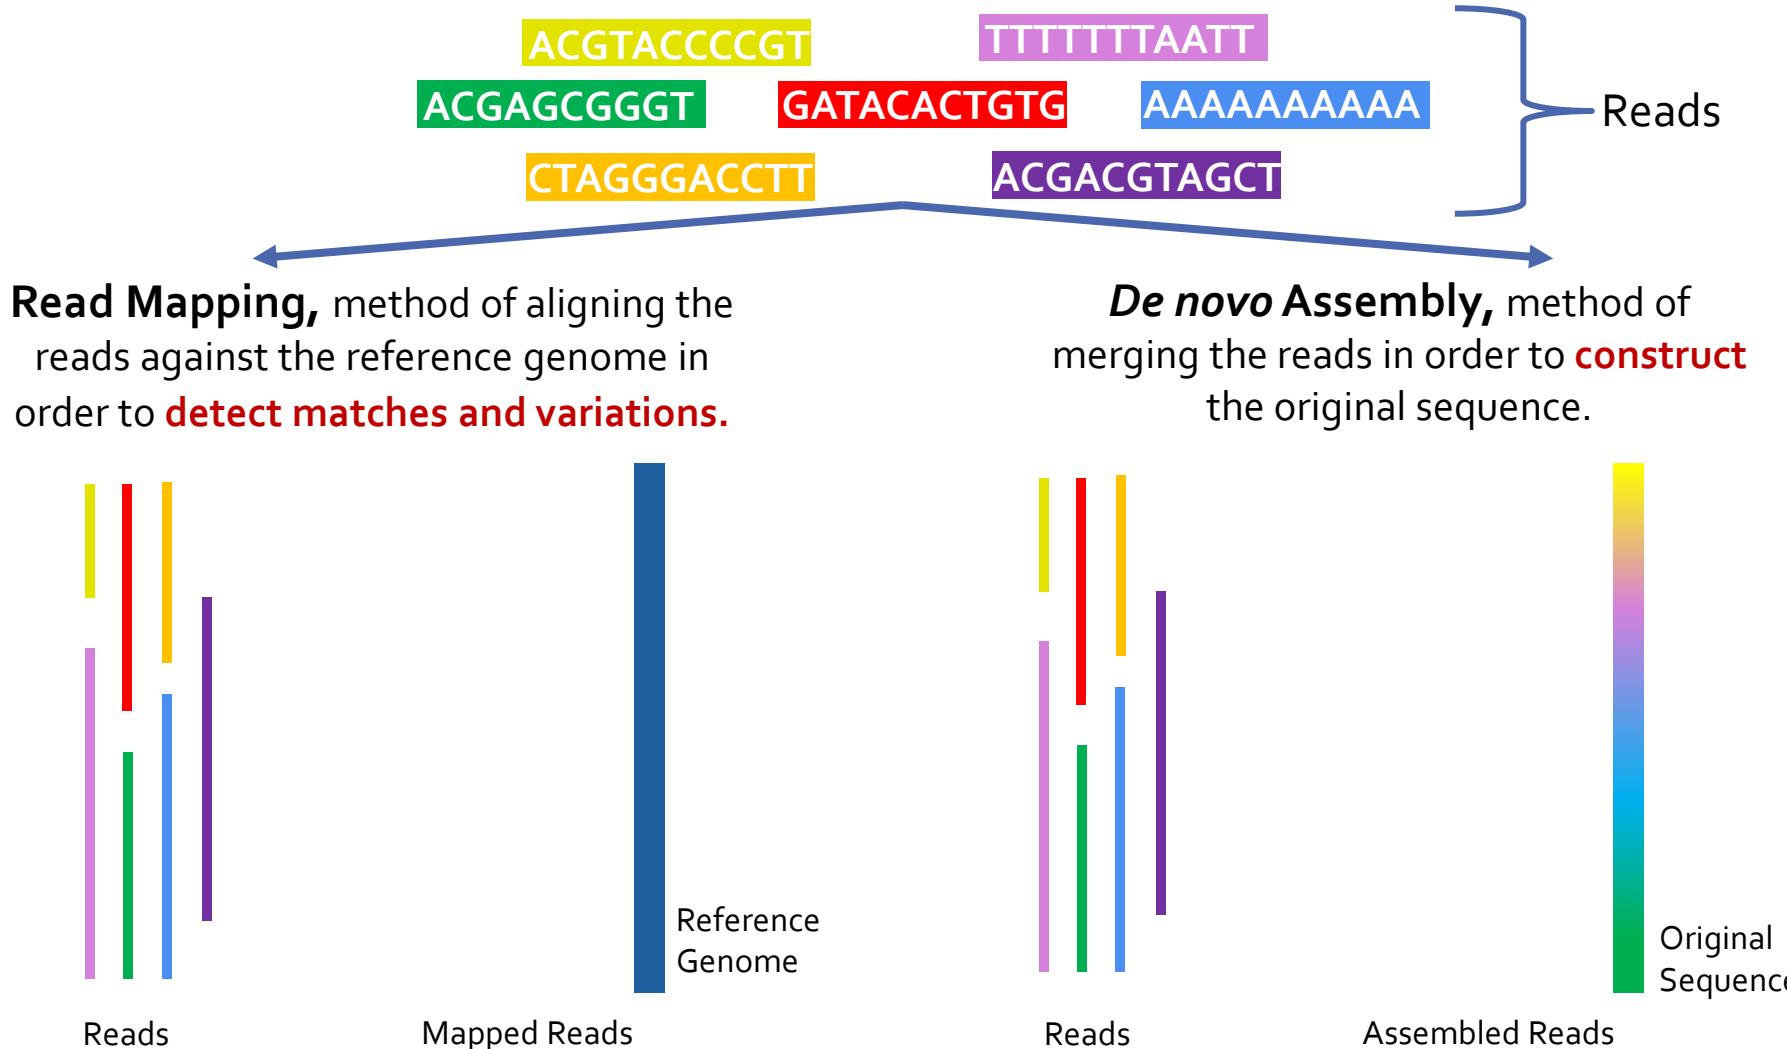
**AACBB Workshop @ ISCA 2022**  
June 18, 2022

# Backup Slides

(BiB Paper)

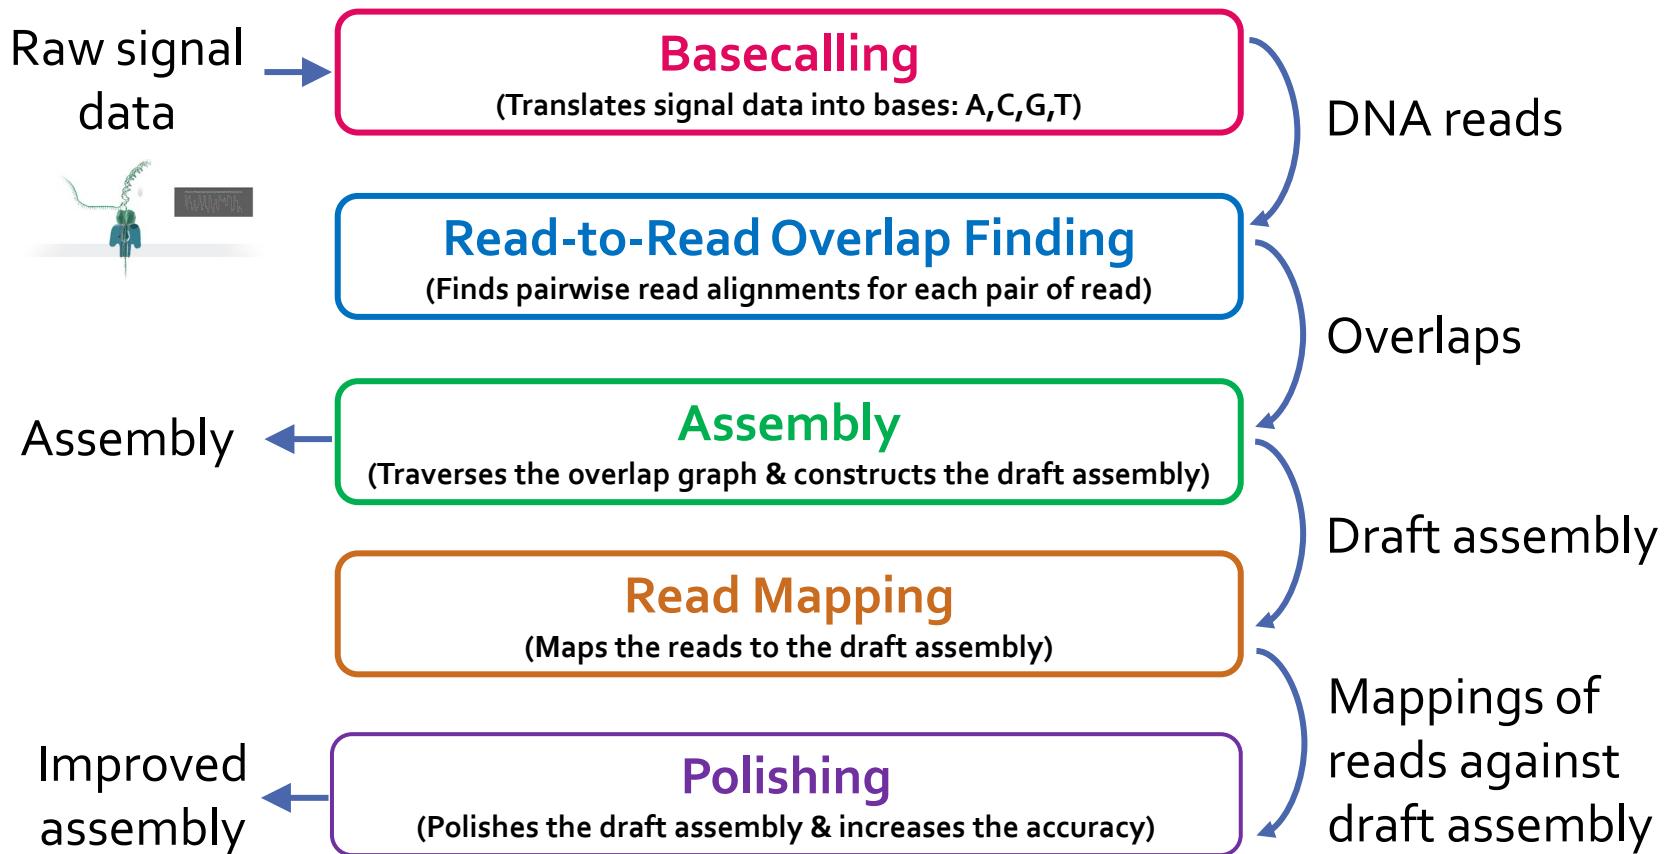
---

# Genome Sequence Analysis



# Genome Assembly Pipeline Using Long Reads

- With the emergence of long read sequencing technologies, *de novo* assembly becomes a promising way of constructing the original genome.



# Our Contributions

---

- Analyze the tools in multiple dimensions: **accuracy**, **performance**, **memory usage**, and **scalability**
- Reveal **new bottlenecks** and **trade-offs**
- First study on **bottleneck analysis** of nanopore sequence analysis pipeline on real machines
- Provide guidelines for **practitioners**
- Provide guidelines for **tool developers**

# Key Findings

---

- Laptops are becoming a popular platform for running genome assembly tools, as the portability of a laptop makes it a good fit for in-field analysis
  - Greater memory constraints
  - Lower computational power
  - Limited battery life
- Memory usage is an important factor that greatly affects the performance and the usability of the tool
  - Data structure choices that increase the memory requirements
  - Algorithms that are not cache-efficient
  - Not keeping memory usage in check with the number of threads
- Scalability of the tool with the number of cores is an important requirement. However, parallelizing the tool can increase the memory usage
  - Not dividing the input data into batches
  - Not limiting the memory usage of each thread
  - Dividing the dataset instead of the computation between simultaneous threads

# Key Findings

**Goal 1:**

**High-performance and low-power**

**Goal 2:**

**Memory-efficient**

**Goal 3:**

**Scalable/highly-parallel**

# Nanopore Sequencing & Tools [BiB 2018]

---

Damla Senol Cali, Jeremie S. Kim, Saugata Ghose, Can Alkan, and Onur Mutlu,  
"Nanopore Sequencing Technology and Tools for Genome Assembly:  
Computational Analysis of the Current State, Bottlenecks and Future  
Directions"

*Briefings in Bioinformatics*, April 2018.

## **Nanopore Sequencing Technology and Tools for Genome Assembly: Computational Analysis of the Current State, Bottlenecks and Future Directions**

**Damla Senol Cali<sup>1,\*</sup>, Jeremie S. Kim<sup>1,3</sup>, Saugata Ghose<sup>1</sup>, Can Alkan<sup>2\*</sup>  
and Onur Mutlu<sup>3,1\*</sup>**

<sup>1</sup>Department of Electrical and Computer Engineering, Carnegie Mellon University, Pittsburgh, PA, USA

<sup>2</sup>Department of Computer Engineering, Bilkent University, Bilkent, Ankara, Turkey

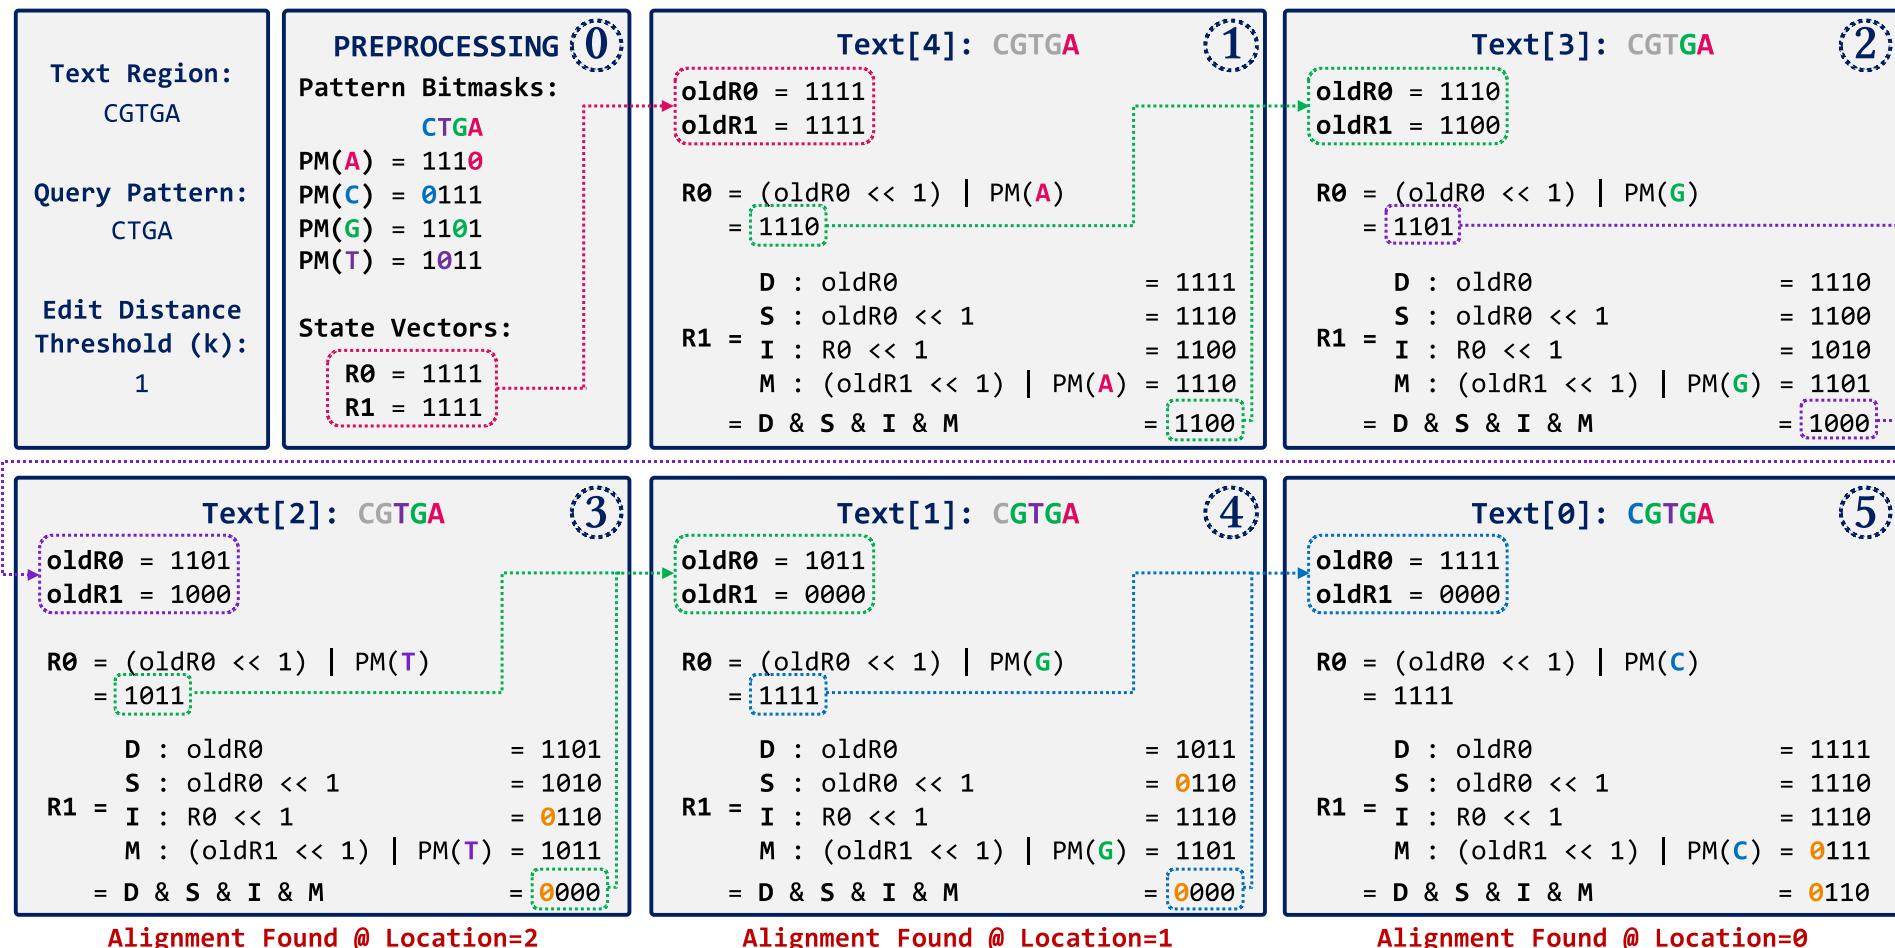
<sup>3</sup>Department of Computer Science, Systems Group, ETH Zürich, Zürich, Switzerland

# Backup Slides

(GenASM)

---

# Example for the Bitap Algorithm



# GenASM Algorithm

---

## □ GenASM-DC Algorithm:

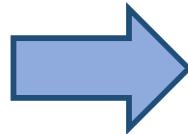
- Modified Bitap for Distance Calculation
- Extended for efficient long read support
- Besides bit-parallelism that Bitap has, extended for parallelism:
  - Loop unrolling
  - Text-level parallelism

## □ GenASM-TB Algorithm:

- Novel Bitap-compatible TraceBack algorithm
- Walks through the intermediate bitvectors (match, deletion, substitution, insertion) generated by GenASM-DC
- Follows a divide-and-conquer approach to decrease the memory footprint

# Loop Unrolling in GenASM-DC

Cycle#	Thread <sub>1</sub> Ro/1/2/..
#1	To-Ro
...	...
#8	To-R <sub>7</sub>
#9	T <sub>1</sub> -Ro
...	...
#16	T <sub>1</sub> -R <sub>7</sub>
#17	T <sub>2</sub> -Ro
...	...
#24	T <sub>2</sub> -R <sub>7</sub>
#25	T <sub>3</sub> -Ro
...	...
#32	T <sub>3</sub> -R <sub>7</sub>



Cycle#	Thread <sub>1</sub> Ro/4	Thread <sub>2</sub> R1/5	Thread <sub>3</sub> R2/6	Thread <sub>4</sub> R3/7
#1	To-Ro	-	-	-
#2	T <sub>1</sub> -Ro	To-R <sub>1</sub>	-	-
#3	T <sub>2</sub> -Ro	T <sub>1</sub> -R <sub>1</sub>	To-R <sub>2</sub>	-
#4	T <sub>3</sub> -Ro	T <sub>2</sub> -R <sub>1</sub>	T <sub>1</sub> -R <sub>2</sub>	To-R <sub>3</sub>
#5	To-R <sub>4</sub>	T <sub>3</sub> -R <sub>1</sub>	T <sub>2</sub> -R <sub>2</sub>	T <sub>1</sub> -R <sub>3</sub>
#6	T <sub>1</sub> -R <sub>4</sub>	To-R <sub>5</sub>	T <sub>3</sub> -R <sub>2</sub>	T <sub>2</sub> -R <sub>3</sub>
#7	T <sub>2</sub> -R <sub>4</sub>	T <sub>1</sub> -R <sub>5</sub>	To-R <sub>6</sub>	T <sub>3</sub> -R <sub>3</sub>
#8	T <sub>3</sub> -R <sub>4</sub>	T <sub>2</sub> -R <sub>5</sub>	T <sub>1</sub> -R <sub>6</sub>	To-R <sub>7</sub>
#9	-	T <sub>3</sub> -R <sub>5</sub>	T <sub>2</sub> -R <sub>6</sub>	T <sub>1</sub> -R <sub>7</sub>
#10	-	-	T <sub>3</sub> -R <sub>6</sub>	T <sub>2</sub> -R <sub>7</sub>
#11	-	-	-	T <sub>3</sub> -R <sub>7</sub>



data *written to memory*  
data *read from memory*

target cell ( $R_d$ )  
cells target cell depends on ( $oldR_d$ ,  $R_{d-1}$ ,  $oldR_{d-1}$ )

# Traceback Example with GenASM-TB

**Deletion Example (Text Location=0)**

**(a)**

Text[0]: C	Text[1]: G	Text[2]: T	Text[3]: G	Text[4]: A
$\begin{cases} R0- : \dots \\ R1-M : 0111 \end{cases}$	$\begin{cases} R0- : \dots \\ R1-D : 1011 \end{cases}$	$\begin{cases} R0-M : 1011 \\ R1- : \dots \end{cases}$	$\begin{cases} R0-M : 1101 \\ R1- : \dots \end{cases}$	$\begin{cases} R0-M : 1110 \\ R1- : \dots \end{cases}$
Match(C) <3,0,1>	Del(-) <2,1,1>	Match(T) <2,2,0>	Match(G) <1,3,0>	Match(A) <0,4,0>

**Substitution Example (Text Location=1)**

**(b)**

Text[1]: G	Text[2]: T	Text[3]: G	Text[4]: A
$\begin{cases} R0- : \dots \\ R1-S : 0110 \end{cases}$	$\begin{cases} R0-M : 1011 \\ R1- : \dots \end{cases}$	$\begin{cases} R0-M : 1101 \\ R1- : \dots \end{cases}$	$\begin{cases} R0-M : 1110 \\ R1- : \dots \end{cases}$
Subs(C) <3,1,1>	Match(T) <2,2,0>	Match(G) <1,3,0>	Match(A) <0,4,0>

**Insertion Example (Text Location=2)**

**(c)**

Text[-]	Text[2]: T	Text[3]: G	Text[4]: A
$\begin{cases} R0- : \dots \\ R1-I : 0110 \end{cases}$	$\begin{cases} R0-M : 1011 \\ R1- : \dots \end{cases}$	$\begin{cases} R0-M : 1101 \\ R1- : \dots \end{cases}$	$\begin{cases} R0-M : 1110 \\ R1- : \dots \end{cases}$
Ins(C) <3,2,1>	Match(T) <2,2,0>	Match(G) <1,3,0>	Match(A) <0,4,0>

# Key Results – Use Case 1

## (1) Read Alignment Step of Read Mapping

- Find the optimal alignment of how reads map to candidate reference regions

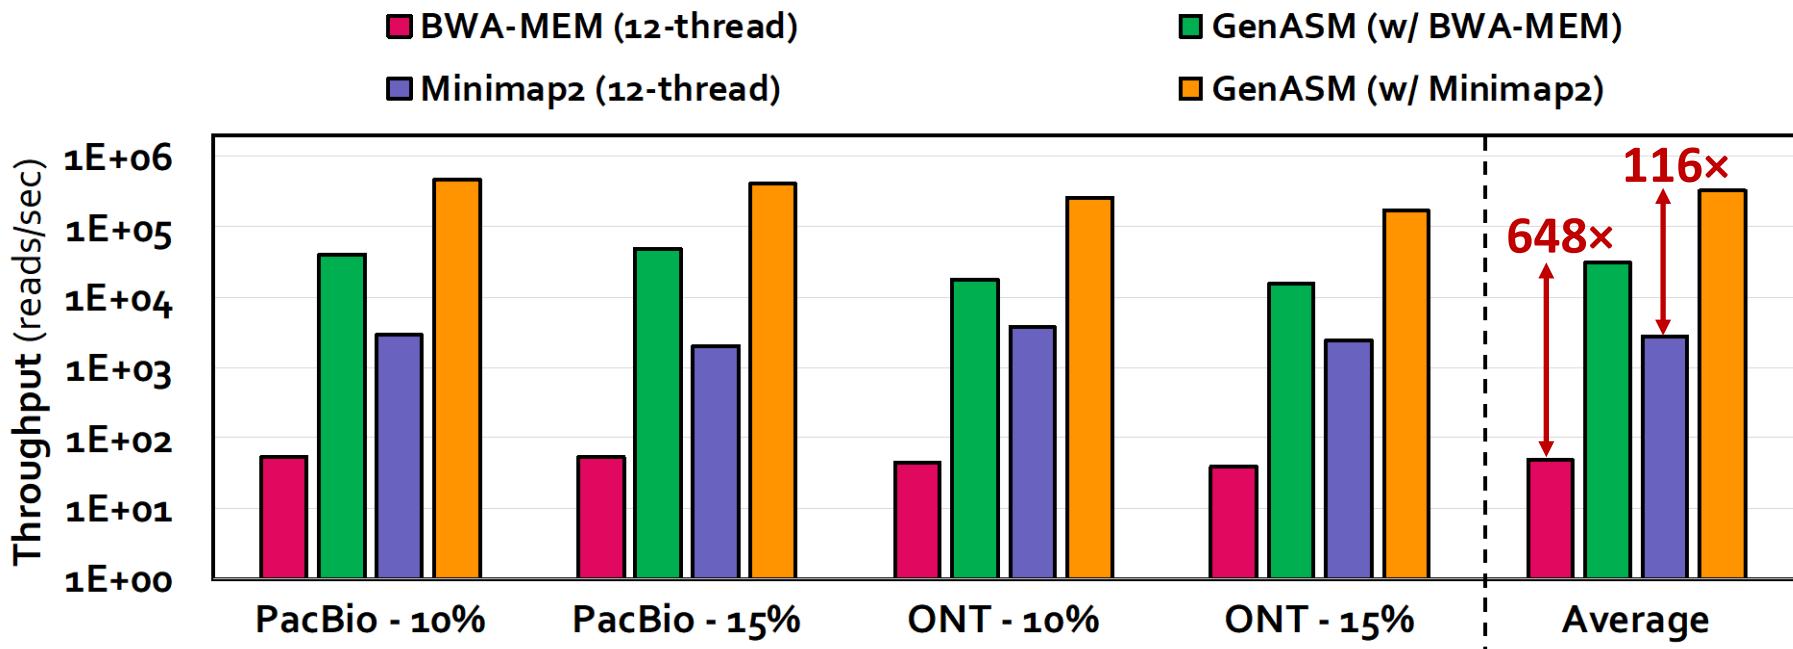
## (2) Pre-Alignment Filtering for Short Reads

- Quickly identify and filter out the unlikely candidate reference regions for each read

## (3) Edit Distance Calculation

- Measure the similarity or distance between two sequences

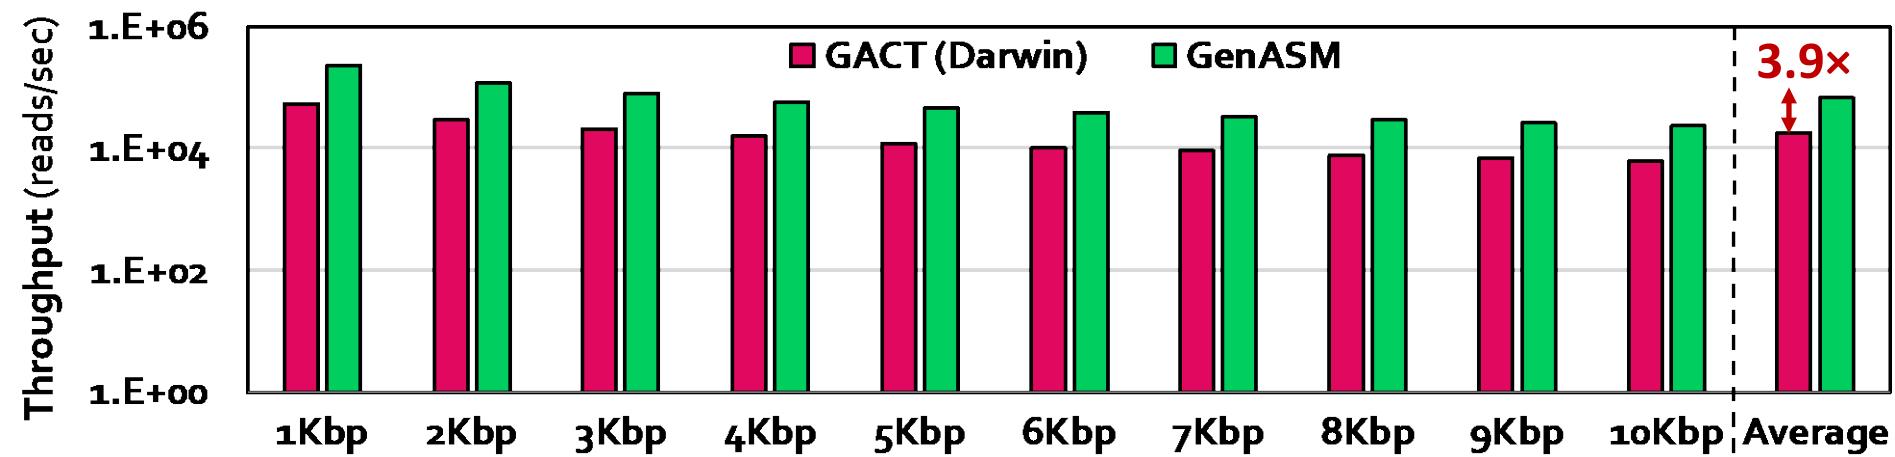
# Key Results – Use Case 1 (Long Reads)



SW

GenASM achieves 648x and 116x speedup over 12-thread runs of BWA-MEM and Minimap2, while reducing power consumption by 34x and 37x

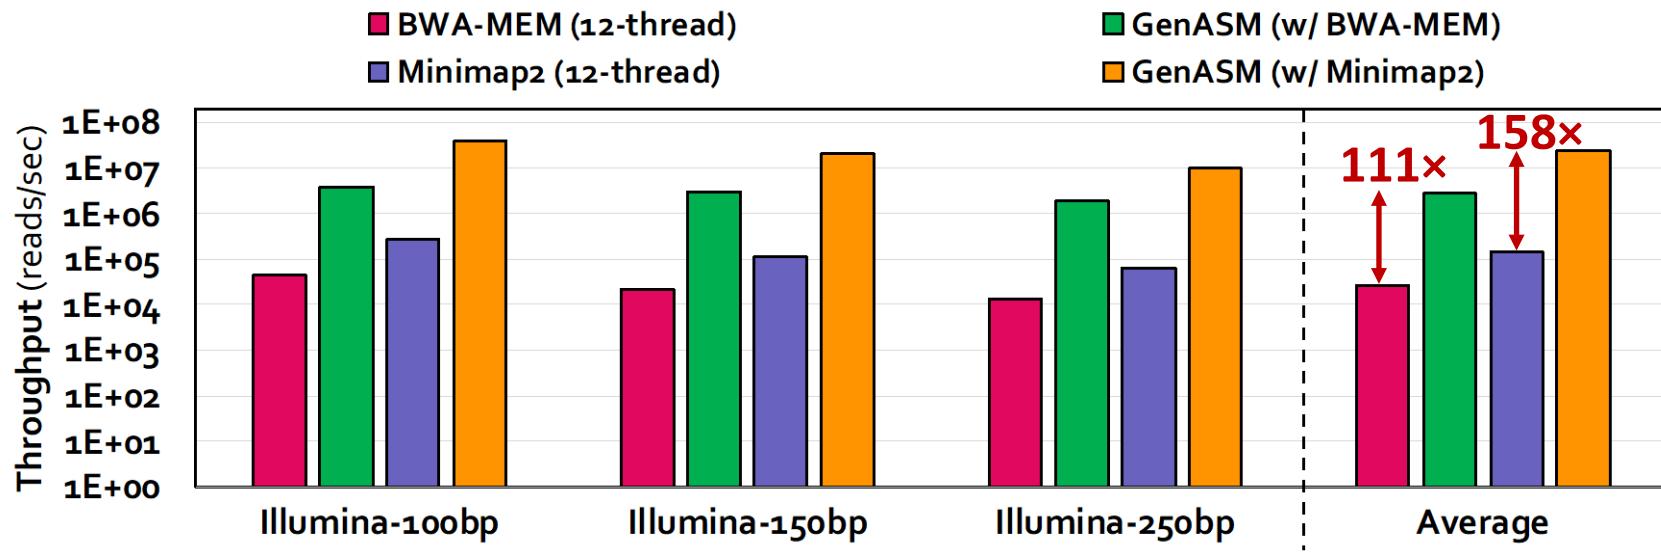
# Key Results – Use Case 1 (Long Reads)



HW

GenASM provides **3.9× better throughput**,  
**6.6× the throughput per unit area**, and  
**10.5× the throughput per unit power**,  
compared to GACT of Darwin

# Key Results – Use Case 1 (Short Reads)



SW

GenASM achieves **111×** and **158×** speedup over  
12-thread runs of BWA-MEM and Minimap2,  
while **reducing power consumption by 33× and 31×**

HW

GenASM provides **1.9× better throughput** and  
**uses 63% less logic area** and **82% less logic power**,  
compared to SillaX of GenAx

# Key Results – Use Case 2

---

## (1) Read Alignment Step of Read Mapping

- Find the optimal alignment of how reads map to candidate reference regions

## (2) Pre-Alignment Filtering for Short Reads

- Quickly identify and filter out the unlikely candidate reference regions for each read

## (3) Edit Distance Calculation

- Measure the similarity or distance between two sequences

# Key Results – Use Case 2

- Compared to Shouji:

- **3.7x** speedup
- **1.7x** less power consumption
- **False accept rate of 0.02%** for GenASM vs. 4% for Shouji
- **False reject rate of 0%** for both GenASM and Shouji

**HW**

GenASM is **more efficient in terms of both speed and power consumption, while significantly improving the accuracy of pre-alignment filtering**

# Key Results – Use Case 3

---

## (1) Read Alignment Step of Read Mapping

- Find the optimal alignment of how reads map to candidate reference regions

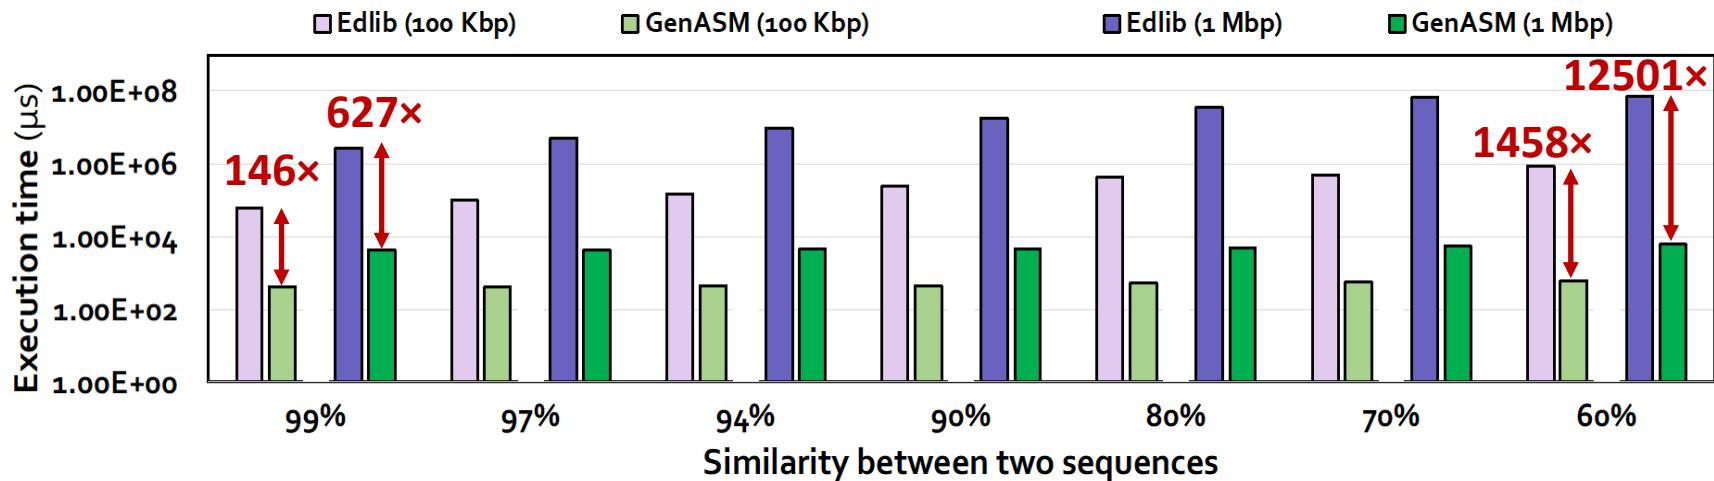
## (2) Pre-Alignment Filtering for Short Reads

- Quickly identify and filter out the unlikely candidate reference regions for each read

## (3) Edit Distance Calculation

- Measure the similarity or distance between two sequences

# Key Results – Use Case 3



SW

GenASM provides **146 – 1458x** and **627 – 12501x speedup**,  
**while reducing power consumption by 548x** and **582x**  
for 100Kbp and 1Mbp sequences, respectively, compared to Edlib

HW

GenASM provides **9.3 – 400x speedup** over ASAP,  
**while consuming 67x less power**

# Sources of Improvement in GenASM

---

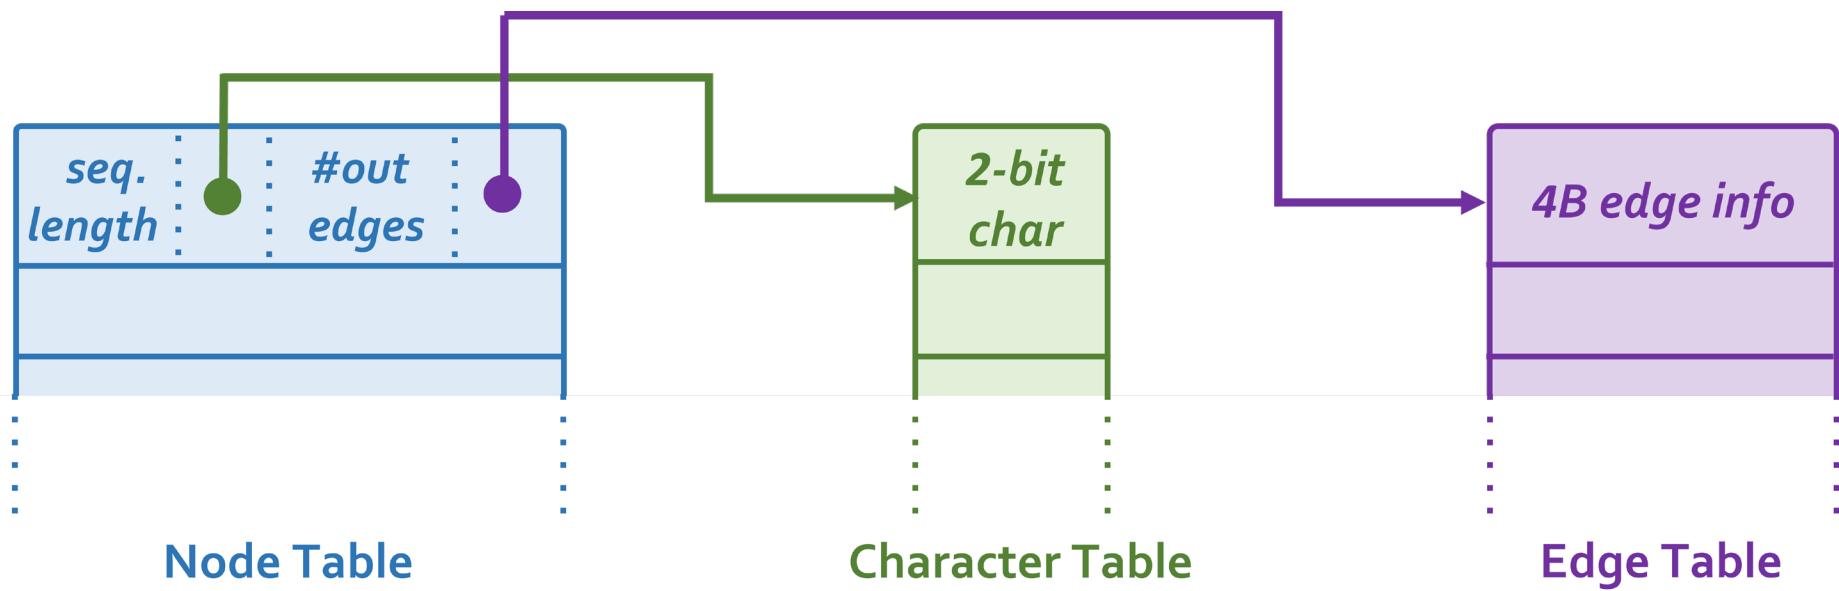
- Very simple computations GenASM performs
- Divide-and-conquer approach we follow, which makes our design efficient for both short and long reads despite their different error profiles
- Very high degree of parallelism obtained with the help of:
  - Specialized compute units, dedicated SRAMs for both GenASM-DC and GenASM-TB, and
  - Vault-level parallelism provided by processing in the logic layer of 3D-stacked memory

# Backup Slides

(SeGram)

---

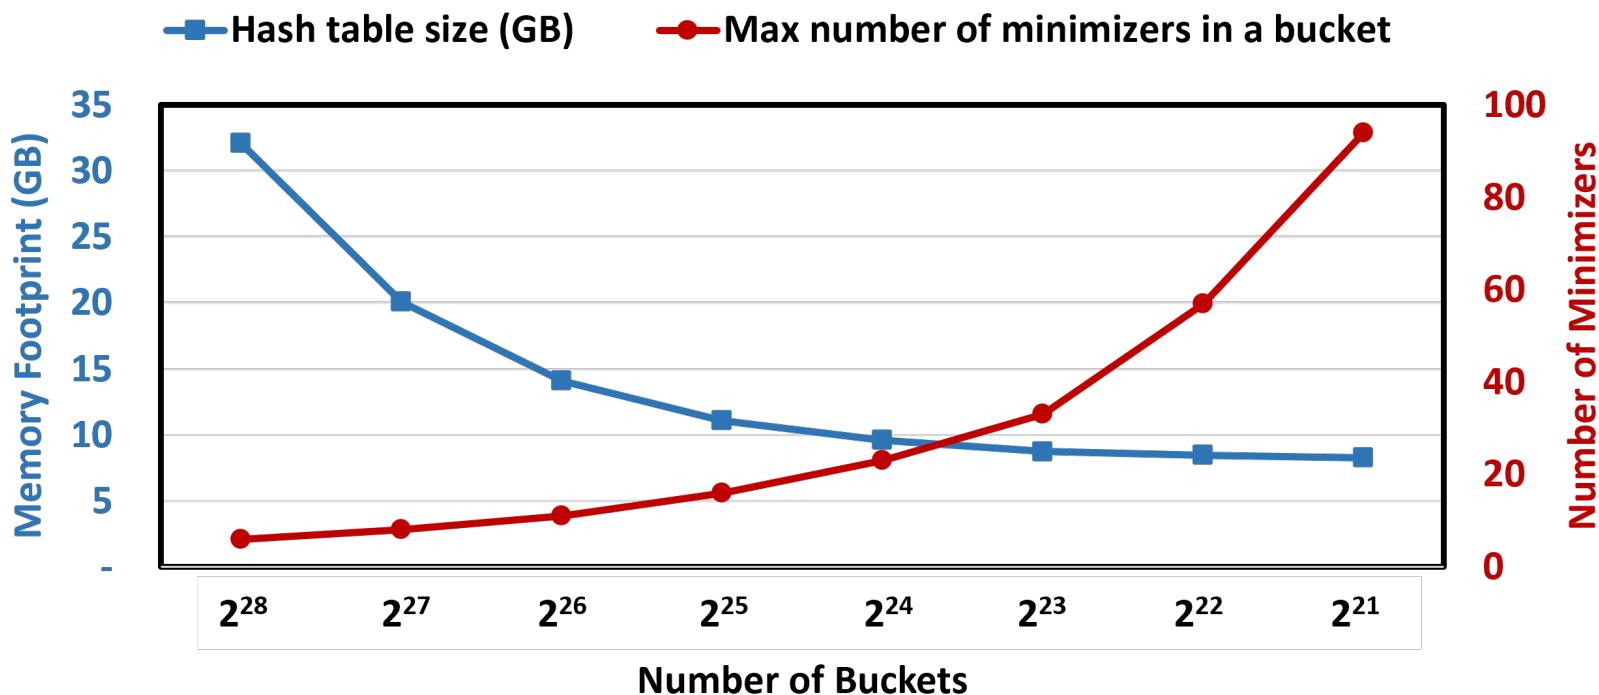
# SeGraM – Graph Structure



# SeGraM – Index Structure



# SeGraM – Selection of #Buckets



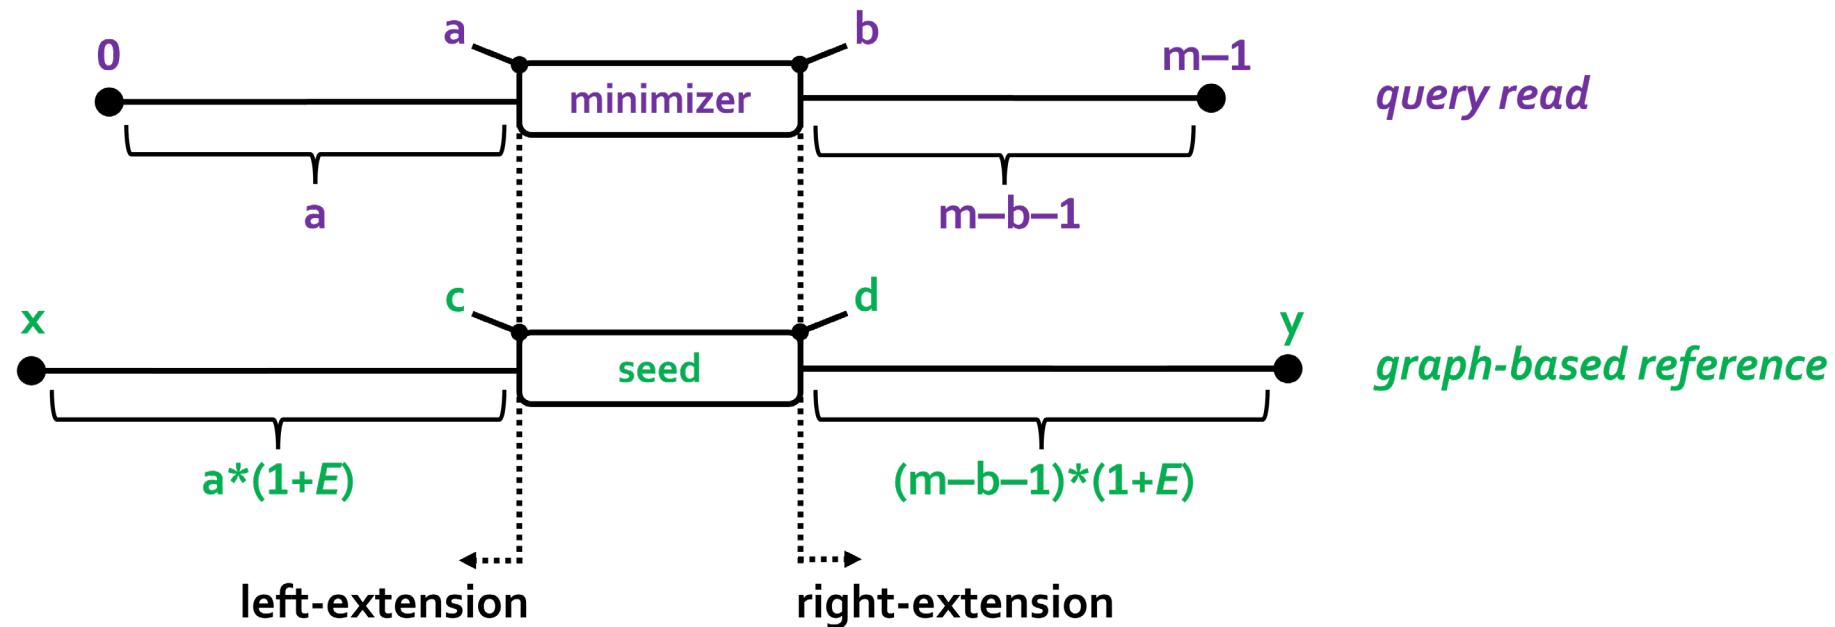
# Minimizers

---

Position	1	2	3	4	5	6	7	...
Sequence	A	G	T	A	G	C	A	...
$k\text{-mer}_1$	A	G	T					
$k\text{-mer}_2$		G	T	A				
$k\text{-mer}_3$			T	A	G			
$k\text{-mer}_4$				A	G	C		
$k\text{-mer}_5$					G	C	A	...

lexicographically  
smallest  $k\text{-mer}$

# MinSeed – Region Calculation



# BitAlign Algorithm

---

## Algorithm 1 BitAlign Algorithm

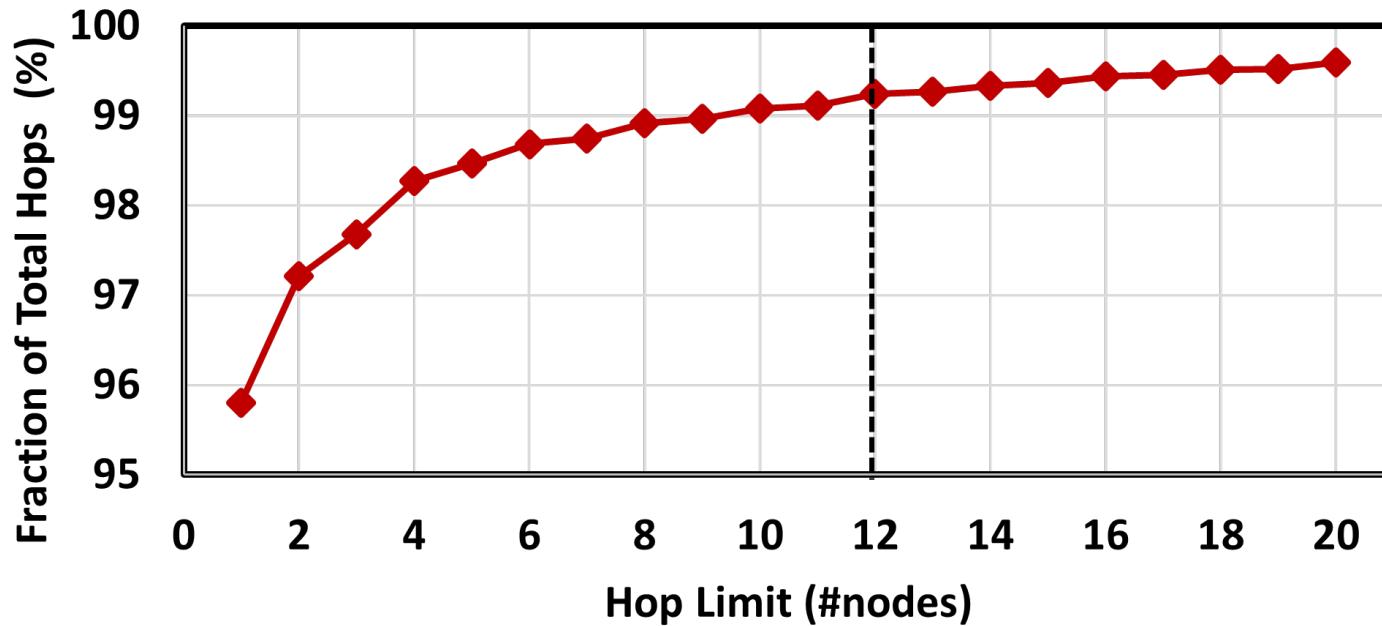
```
Inputs: linearized and topologically sorted subgraph (reference),
       query-read (pattern), k (edit distance threshold)
Outputs: editDist (minimum edit distance), CIGARstr (traceback output)

1: n ← length of linearized reference subgraph
2: m ← length of query read
3: PM ← genPatternBitmasks(query-read)           ▷ pre-process the query read
4:
5: allR[n][d] ← 111...111      ▷ init R[d] bitvectors for all characters with 1s
6:
7: for i in (n-1):-1:0 do          ▷ iterate over each subgraph node
8:   curChar ← subgraph-nodes[i].char
9:   curPM ← PM[curChar]           ▷ retrieve the pattern bitmask
10:
11:  R0 ← 111...111               ▷ status bitvector for exact match
12:  for j in subgraph-nodes[i].successors do
13:    R0 ← ((R[j][0]<<1) | curPM) & R0      ▷ exact match calculation
14:  allR[i][0] ← R0
15:
16:  for d in 1:k do
17:    I ← (allR[i][d-1]<<1)                  ▷ insertion
18:    Rd ← I                                ▷ status bitvector for d errors
19:    for j in subgraph-nodes[i].successors do
20:      D ← allR[j][d-1]                      ▷ deletion
21:      S ← allR[j][d-1]<<1                 ▷ substitution
22:      M ← (allR[j][d]<<1) | curPM        ▷ match
23:      Rd ← D & S & M & Rd
24:    allR[i][d] ← Rd
25: <editDist, CIGARstr> ← traceback(allR, subgraph, query-read)
```

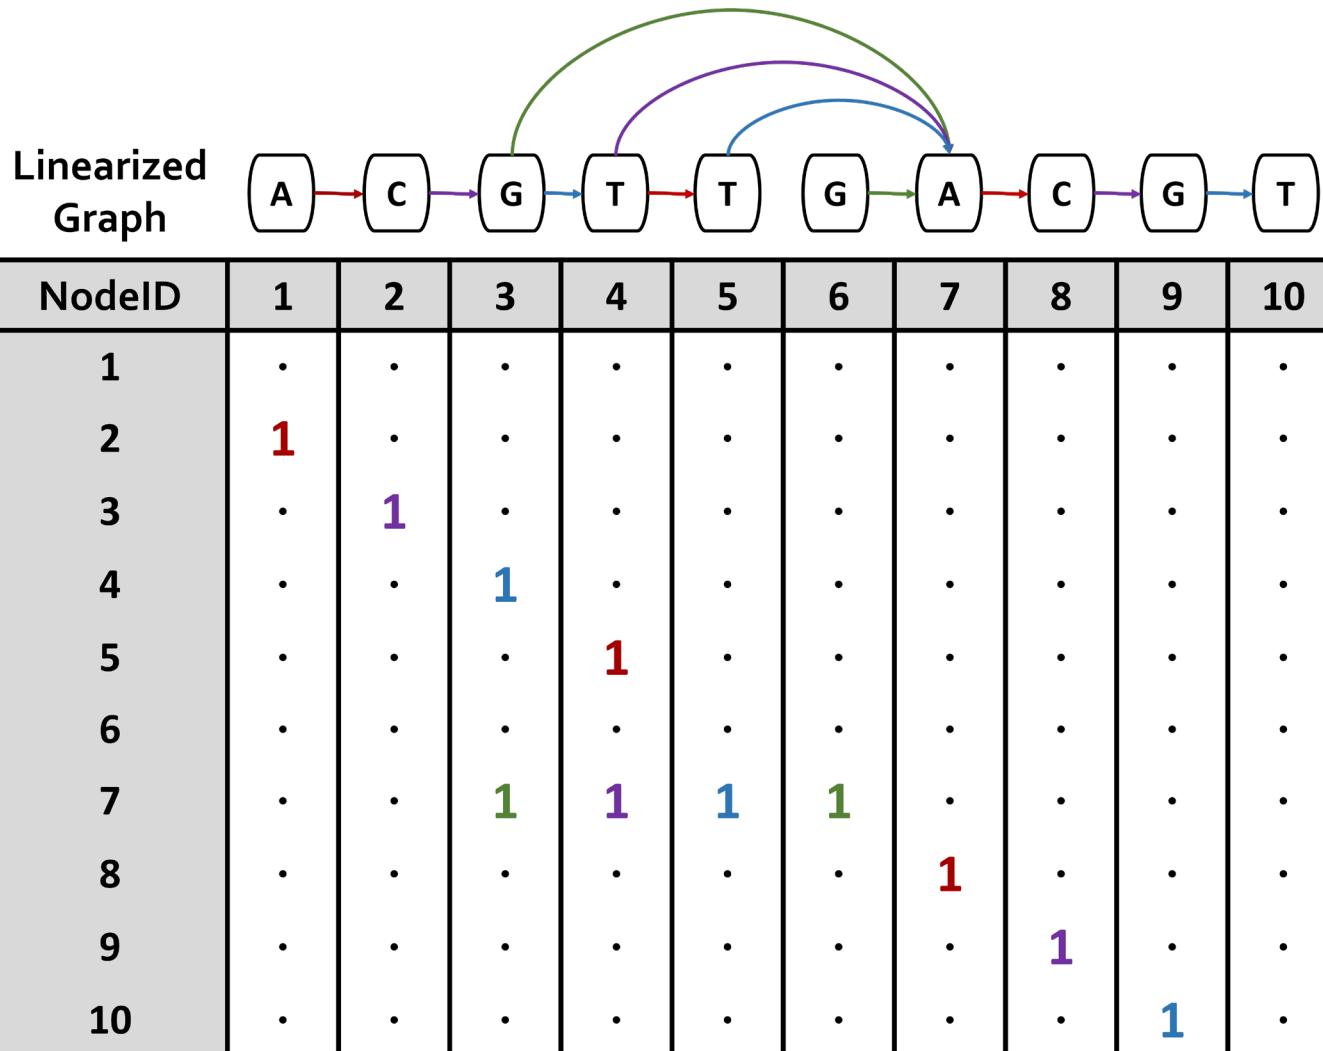
---

# BitAlign – Hop Length Selection

---



# BitAlign – HopBits



# Sources of Improvement

---

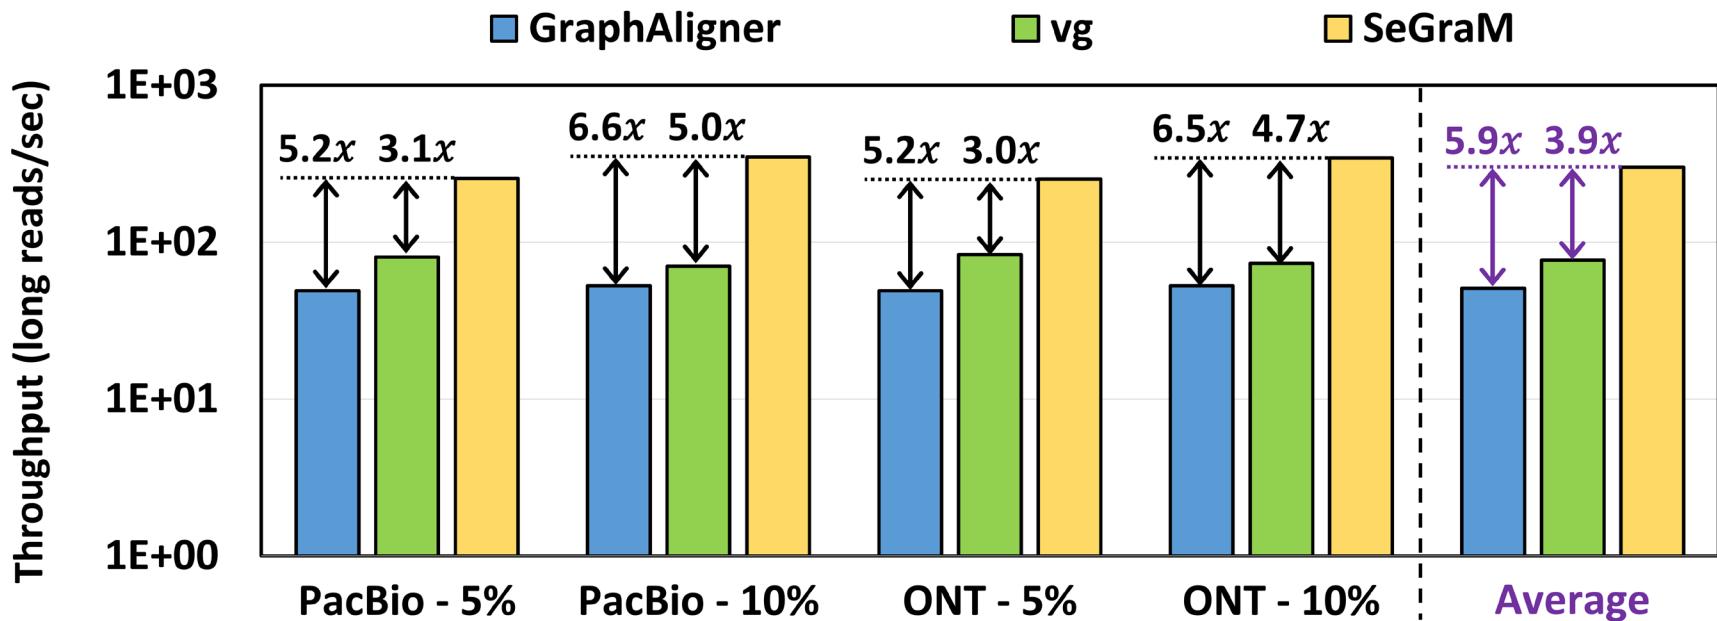
- **Co-design approach for both seeding and alignment:**
  - Efficient and hardware-friendly algorithms for seeding and for alignment
  - Eliminating the data transfer bottleneck between the seeding and alignment steps of the genome sequence analysis pipeline, by placing their individual accelerators (MinSeed and BitAlign) adjacent to each other
  - Pipelining of the two accelerators within a SeGram accelerator, which allows us to completely hide the latency of MinSeed
- **Overcoming the high cache miss rates** observed from the baseline tools by carefully designing and sizing the on-chip scratchpads and the hop queue registers and matching the rate of computation for the logic units with memory bandwidth and memory capacity

# Sources of Improvement (cont'd.)

---

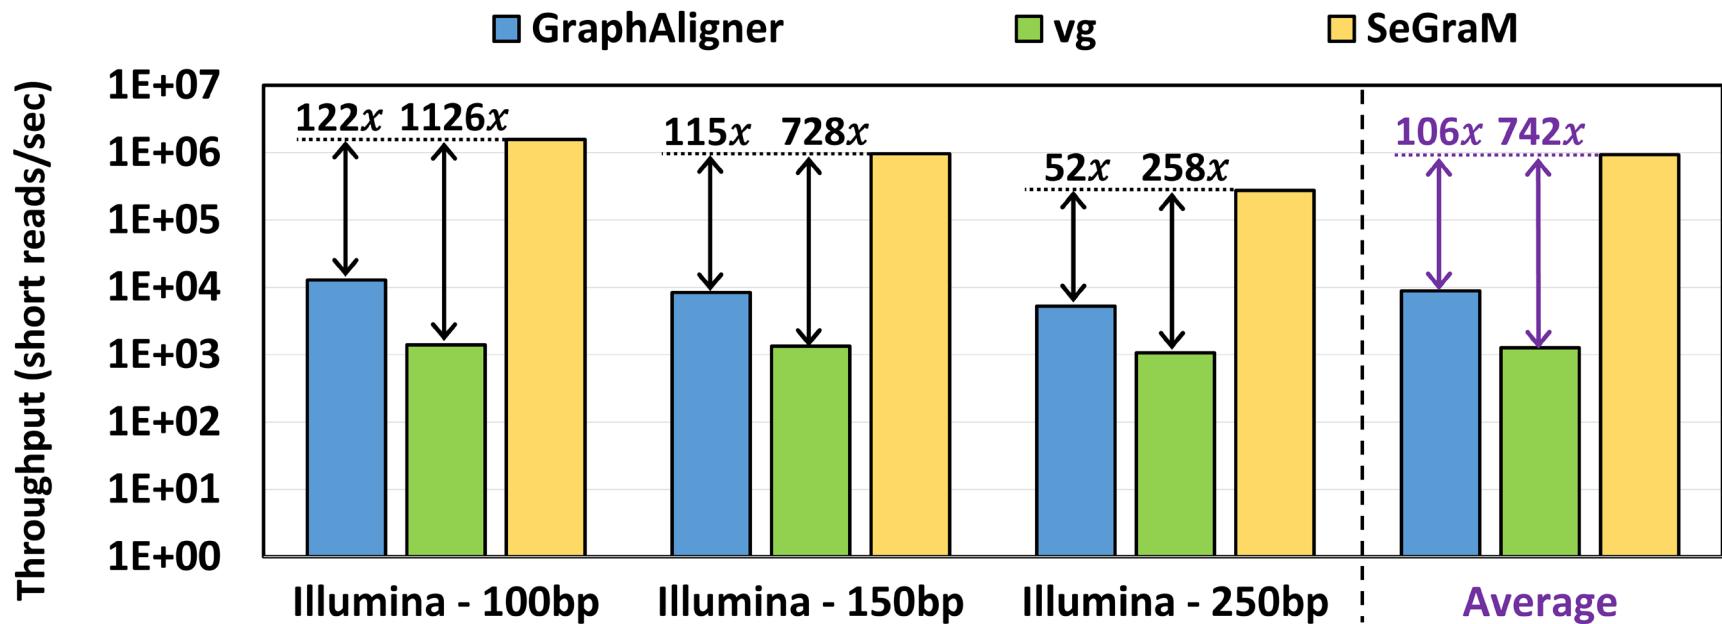
- **Addressing the DRAM latency bottleneck** by taking advantage of the natural channel subdivision exposed by HBM and eliminating any inter-accelerator interference-related latency in the memory system
- **Scaling linearly across three dimensions:**
  - Within a single BitAlign accelerator, by incorporating processing elements (*i.e., iteration-level parallelism*),
  - Executing multiple seeds in parallel by using pipelined execution with the help of our double buffering approach (*i.e., seed-level parallelism*), and
  - Processing multiple reads concurrently without introducing inter-accelerator memory interference with the help of multiple HBM stacks that each contain the same content (*i.e., read-level parallelism*)

# Key Results – SeGraM with Long Reads



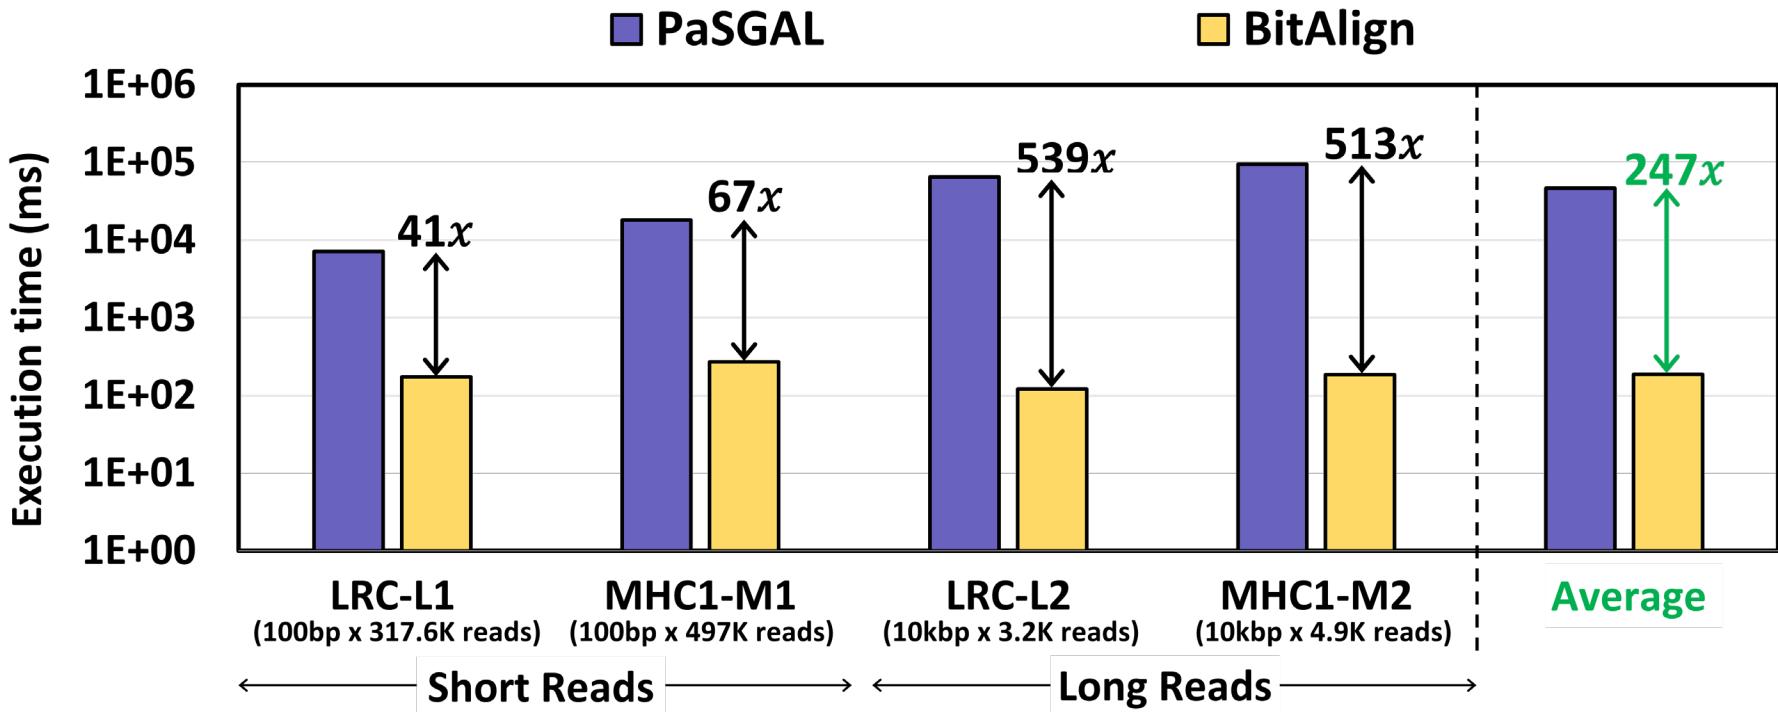
SeGraM provides **5.9x** and **3.9x** throughput improvement  
over GraphAligner and vg,  
while reducing the power consumption by **4.1x** and **4.4x**

# Key Results – SeGraM with Short Reads



SeGraM provides **106x and 742x throughput improvement** over GraphAligner and vg,  
while **reducing the power consumption by 3.0x and 3.2x**

# Key Results – BitAlign (S2G Alignment)



BitAlign provides  **$41\times$ - $539\times$  speedup** over PaSGAL

# Key Results – BitAlign (S2S Alignment)

---

- BitAlign can be used for both sequence-to-sequence alignment and sequence-to-graph alignment
  - The cost of more functionality: **Extra hop queue registers** in BitAlign
  - **We do *not* sacrifice any performance**
- **For long reads (over GACT of Darwin and GenASM):**
  - 4.8× and 1.2× throughput improvement,
  - 2.7× and 7.5× higher power consumption, and
  - 1.5× and 2.6× higher area overhead
- **For short reads (over SillaX of GenAx and GenASM):**
  - 2.4× and 1.3× throughput improvement

# Backup Slides

(BitMAc)

---

# BitMAC: FPGA-based GenASM

## Our Goal:

Map GenASM accelerators to an FPGA with HBM2,  
where **HBM2 offers high memory bandwidth** and

**FPGA resources offer high parallelism** by instantiating  
multiple copies of the GenASM accelerators

- ❑ **Re-modified GenASM algorithms** for a better mapping to the FPGA resources
- ❑ **Intra-level parallelism** by instantiating multiple processing elements (PEs) for the DC execution
- ❑ **Inter-level parallelism** by running multiple independent GenASM executions in parallel

# Key Findings

---

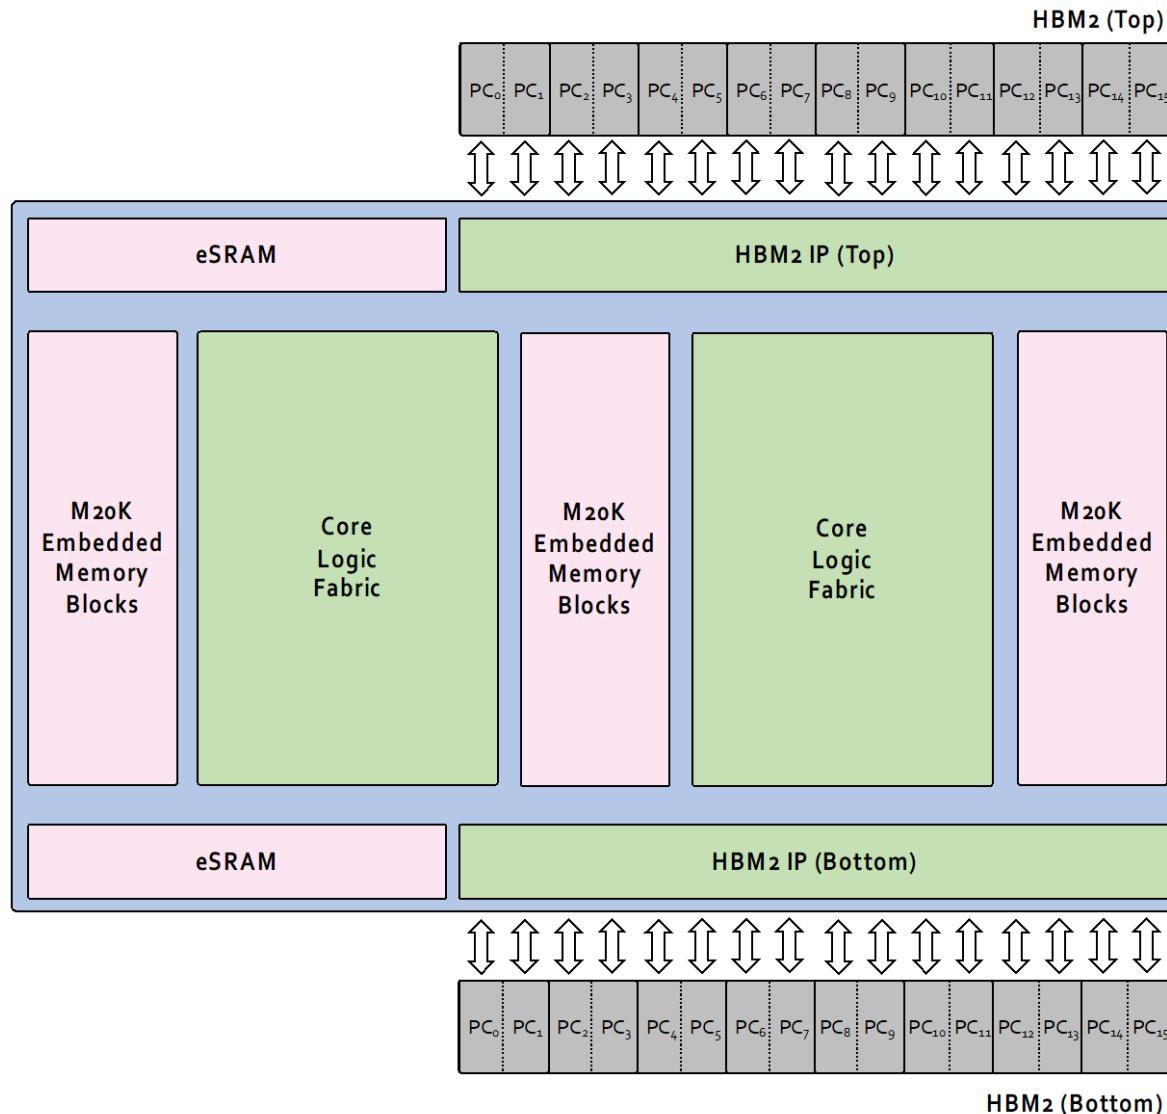
- Based on the FPGA resources, **the complete BitMAC design:**
  - 4 BitMAC accelerators connected to each pseudo-channel (128 in total)
  - Each BitMAC accelerator contains a DC accelerator with 16 PEs, a TB accelerator, an FSM, and 13.2KB of M20Ks
  - Clocked at 200MHz
  
- BitMAC provides:
  - **136x – 761x speedup** over the state-of-the-art CPU baselines
  - **6.8x – 19.4x speedup** over the state-of-the-art GPU baseline

# Key Findings (cont'd.)

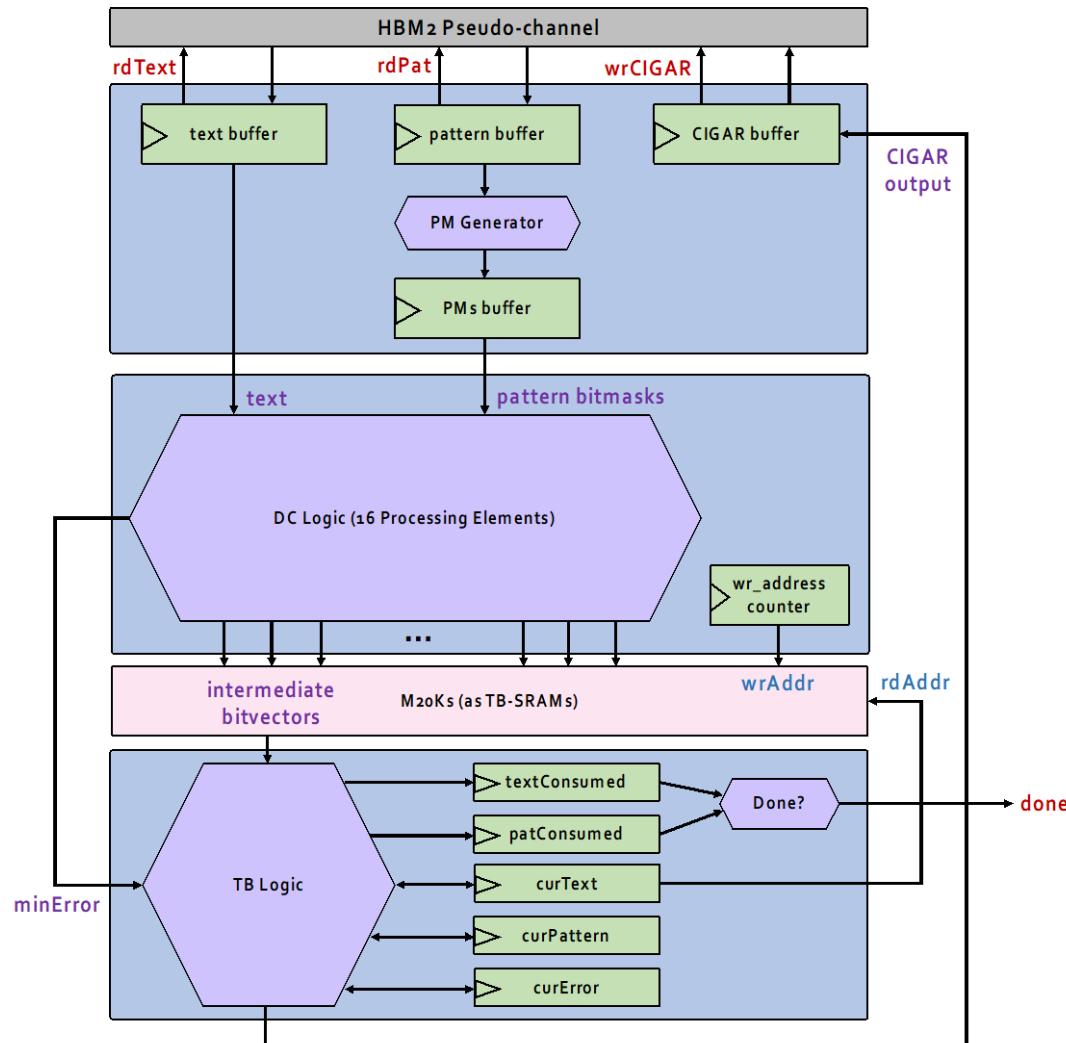
---

- BitMAC has:
  - **64% logic utilization** and **90% on-chip memory utilization**
  - Total power consumption of **48.9W**, where **59%** accounts for the **M20Ks**
- **Bottlenecked by the amount of on-chip memory** (i.e., M20Ks)
- **Cannot saturate the high bandwidth** that multiple HBM2 stacks on the FPGA provide
- Need (1) **algorithm-level modifications** to decrease the amount of data that need to be stored in M20Ks, and (2) **newer FPGA chips that provide a higher amount of on-chip memory capacity**

# Intel Stratix 10 MX



# BitMAC Design



# BitMAC – Results

---

Component	Dynamic On-Chip Power Dissipation	Total On-Chip Power Dissipation
DC Logic (16 PEs)	128.57 mW	
TB Logic	10.24 mW	
FSM Logic	3.15 mW	
M2oKs	211.61 mW	
Other	15.72 mW	
<b>Total – 1 BitMAC Accelerator</b>	<b>369.29 mW (0.4 W)</b>	<b>6043.24 mW (6.0 W)</b>
<b>Total – 32 BitMAC Accelerators (1 per each pseudo-channel)</b>	<b>11569.92 mW (11.6 W)</b>	<b>17234.67 mW (17.2 W)</b>
<b>Total – 128 BitMAC Accelerators (4 per each pseudo-channel)</b>	<b>43042.90 mW (43 W)</b>	<b>48935.65 mW (48.9 W)</b>

# BitMAC – Results

---

Configuration	Logic Utilization	M2oK	eSRAM	DSP
1 BitMAC Accelerator	0.5%	0.7%	0%	0%
32 BitMAC Accelerators (1 per each pseudo-channel)	17.7%	22.4%	0%	0%
128 BitMAC Accelerators (4 per each pseudo-channel)	64.3%	89.7%	0%	0%